AI PRODUCT

PORTFOLIO.

AI Product Manager, trained as an architect

Hello, I'm Wenyuan.

AI Product Manager, trained as an architect

I'm Wenyuan Li (Alnt Med). Trained as an architect, now building AI products — I run a department agent platform at XTOOL and build my own. Architecture taught me to turn scattered needs into a stable system; AI product work is where that method meets real scenarios.

EMAIL lee.menuong@gmail.com

GITHUB github.com/leemenuong-prog

SITE alnt-med.netlify.app

EDUCATION

2025 –	Shenzhen University · M.Arch, architecture × AI — the pivot to AI product
2020–2025	Guangzhou University · B.Arch

EXPERIENCE

2026	XTOOL · AIPM intern · Co-work agent platform, zero to one (65 people across 4 departments, 80%+ completion)
2025	Puwei Tech · AI product intern (vertical RAG agents)
2025	UABB Bi-city Biennale · multimodal AIGC lead · sole student representative
2025	Architectural Journal (T1) · second author, AIGC spatial optimisation

AWARDS

Challenge Cup · Grand Prize	TOP 1%
Lanqiao Cup, national software competition · National First Prize	TOP 1%
Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up	TOP 3%
CUHK Enterprise Agent Hackathon · Runner-up	TOP 5%
NCDA National Digital Design Award · First Prize	

SKILLS

PRODUCT	AI coding (Claude Code / Cursor) · prototyping (Figma / Axure / ClaudeDesign) · PRD
DATA	SQL · Python · A/B testing · AI eval sets · Apify · Excel pivots
ORCHESTRATION	Coze · n8n
MULTIMODAL	Tripo 3D · ComfyUI · Jimeng · Midjourney · UE5

alnt-med.netlify.app

Scan or click for the live site ↗

INDEX.

AI PRODUCT × ARCHITECTURE

Pears 01

AI Product · 2026

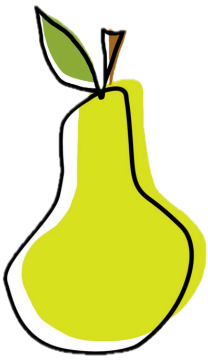

Co-work 02

AI Platform · 2026

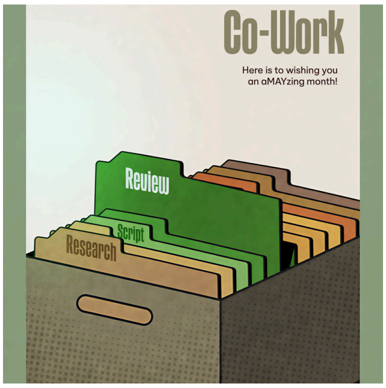

Yijian 03

AI Product · 2026

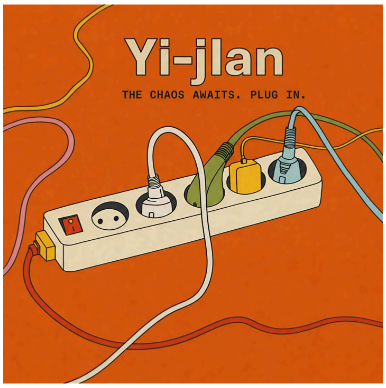

Meco 04

AI Product · 2026

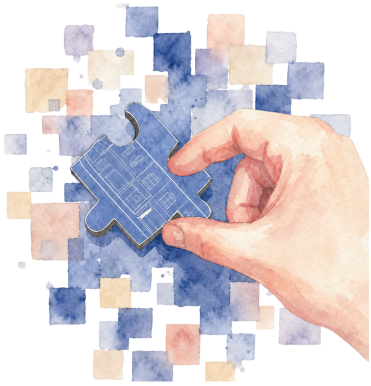

UABB 05

AIGC Pipeline · 2025

Architecture 06

ARCHITECTURE · 2023–2025

PEARS.

Private agent platform · Work distillation · Workflow agents

The Short Version

— AI has leapt ahead these past two years, but I kept seeing the same gap: teams still dragging through repetitive work that an agent could take off their hands.

— I wanted to move the bar for building an agent from "explain the requirement to an AI" to "just do the task once and let it watch."

— **Solo,** I built a private-domain agent platform: it watches you work, distills the trace into a PRD you can read and edit, then hands it to an AI coding agent that builds a workflow agent of your own — zero to launch.

— It took 2nd runner-up at the Tsinghua AI Agent collegiate hackathon, and drew some of the most enterprise interest on site.

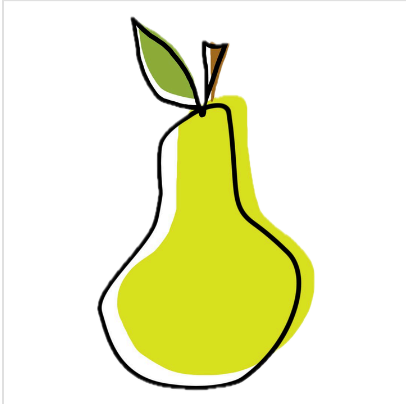

Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up

4+	36/85	0→1	TOP 3%
ENTERPRISE LEADS ON SITE	PEER VOTE · NO.1	SOLO · ZERO TO LAUNCH	TSINGHUA · 2ND RUNNER-UP

THE LANDING PAGE

HAND THE BUSYWORK TO A PEAR.

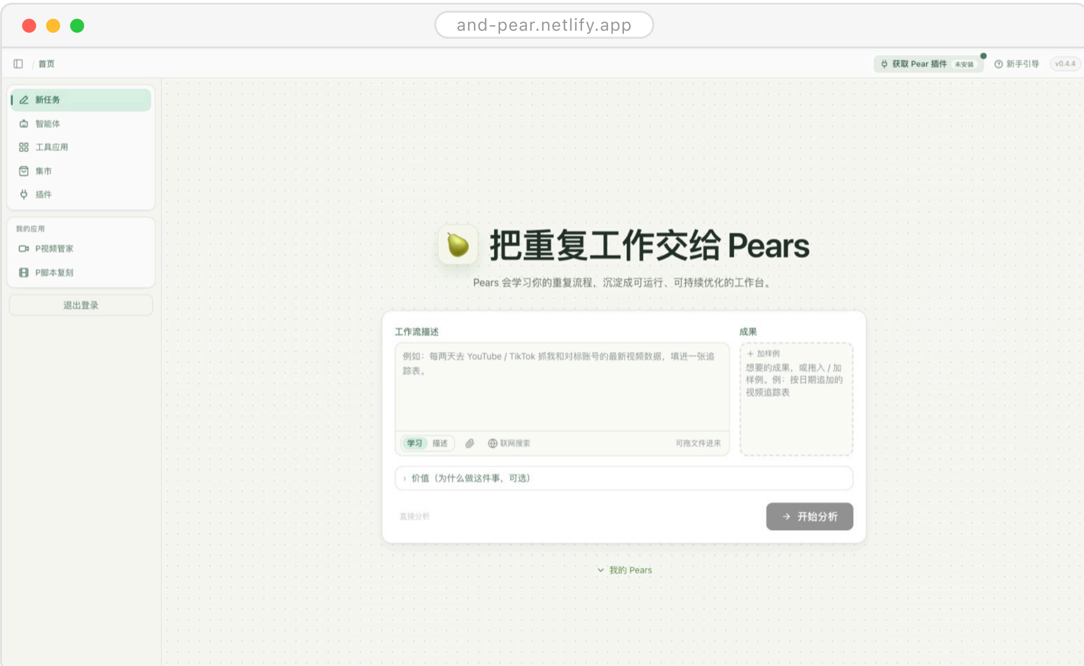

Home · one brief, one sample — learning starts

Watch, distill, build.

Hi, I'm Pears 🍐 — show me your work once, and I turn it into a workflow agent of your own.

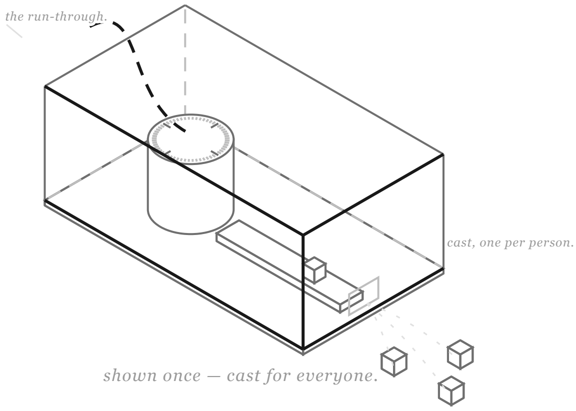

SCENE 01 · THE BRIEF

Why

- AI coding brought the bar for writing software way down, yet the front-line colleagues who understand the business best still don't have an "AI employee" of their own.
- **The catch:** the people who know the work are usually the worst at describing it — ask them to write prompts or draw flowcharts and you've just raised a new barrier.
- But everyone can do the job they do every day. Treat "doing it once" as the spec, and the barrier is gone.

Key Contributions

- **Ran the whole thing solo,** product definition to launch — no team to divide the work or fall back on.
- Designed the observe → distill → generate path: the trace never turns straight into code; it first lands as a PRD you can read and edit, and how far the automation goes is your call.
- Privacy was fixed on day one: it records only when you say so, and stops the instant you stop the task.
- 2nd runner-up at the live roadshow, among the most enterprise-courted projects there (4+ companies), and first in the contestants' own popular vote (36/85).

How It Works

01 Pears
02
03
04
05
06

- **Align:** a short clarifying chat and a goal doc get it clear on what this task should produce.
- Demonstrate: you do the work once in the browser as usual, and the extension records the key moves — what you clicked, switched to, typed.
- **Rebuild:** the trace isn't replayed verbatim; RAG reworks it into a workflow that actually runs.
- **Share:** finished capabilities go into an app market — no teaching from scratch, just swap in an app.
- **Boundaries:** it records only after you explicitly start a task, cuts off the moment you stop, and logs only the key events.

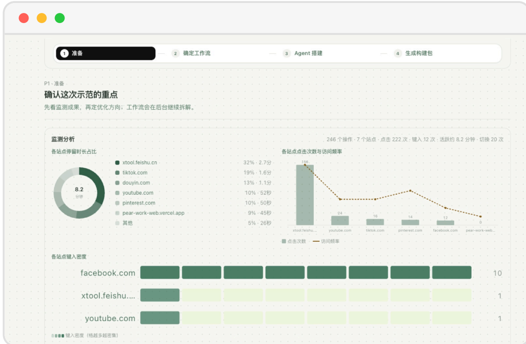

① Watch · clicks, switches, keystrokes

② Assemble · each step, your call

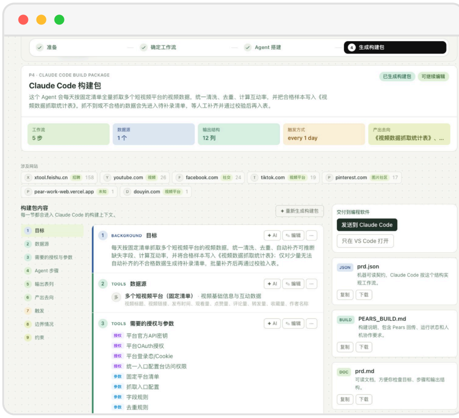

③ Build pack · a PRD you can read and edit

Impact

01 Pears
02
03
04
05
06

- Vibe coding starts from what you say; Pears starts from what you do — and the second never asks you to be good at stating requirements.
- The agent backlog one dev team could never clear becomes something every colleague can clear on their own.
- **A week after launch,** Codex shipped Record & Replay too — outside confirmation that the need here is real.

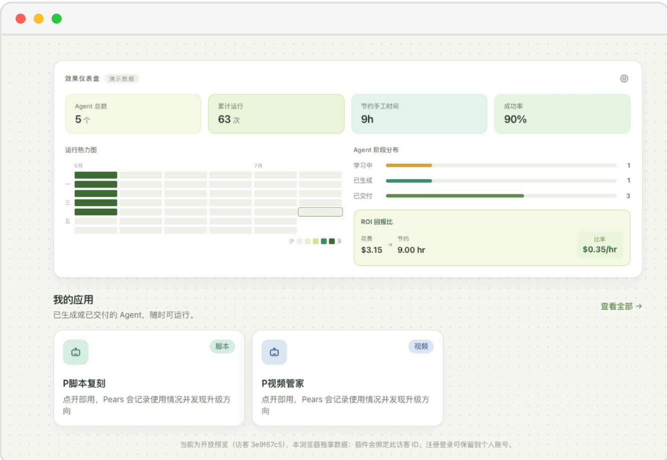

Dashboard · runs, hours saved, ROI

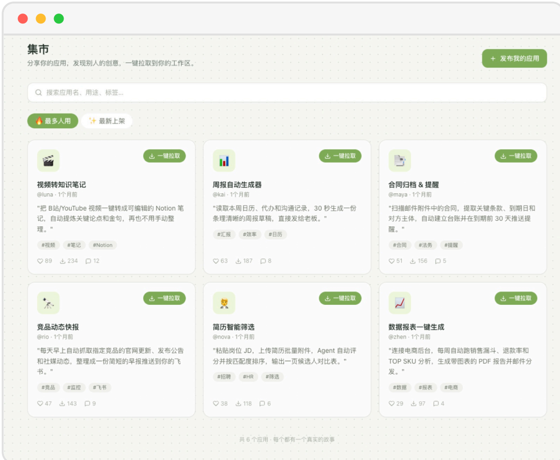

Market · pull a working app into your space

Watch the pitch.

WATCH THE ROADSHOW FILM ↗

STACK Browser extension · PRD distillation · AI coding agents

PROJECT 02 · AI PLATFORM · 2026

Co-work.

Agent platform · RAG · Dashboard · Feishu

The Short Version

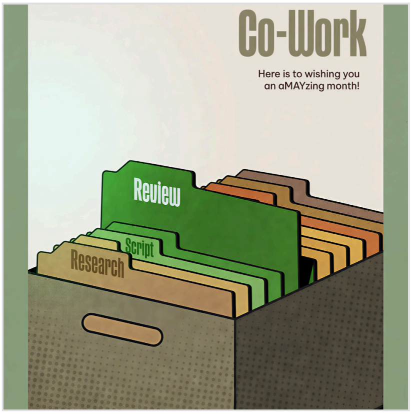

- **Interning at XTOOL,** the department's output was dragged by two things: manual repetition and human-run processes — and worse, know-how left with whoever quit.
- My job was to turn that manual work into agent workflows, and keep everyone's know-how in the platform.
- **Team by team,** I built the department's agent platform zero to one: 8 tools across operations, product and content, iterated on analytics.
- It now serves 65 people across 4 departments at 80%+ completion; by the platform's ROI, every $1 in saves 0.39 work-hours.

65	8	80%+	0.39h
PEOPLE · 4 DEPTS	PRODUCTION TOOLS	COMPLETION RATE	SAVED PER $1

THE LANDING PAGE

Click once — the agents do the rest.

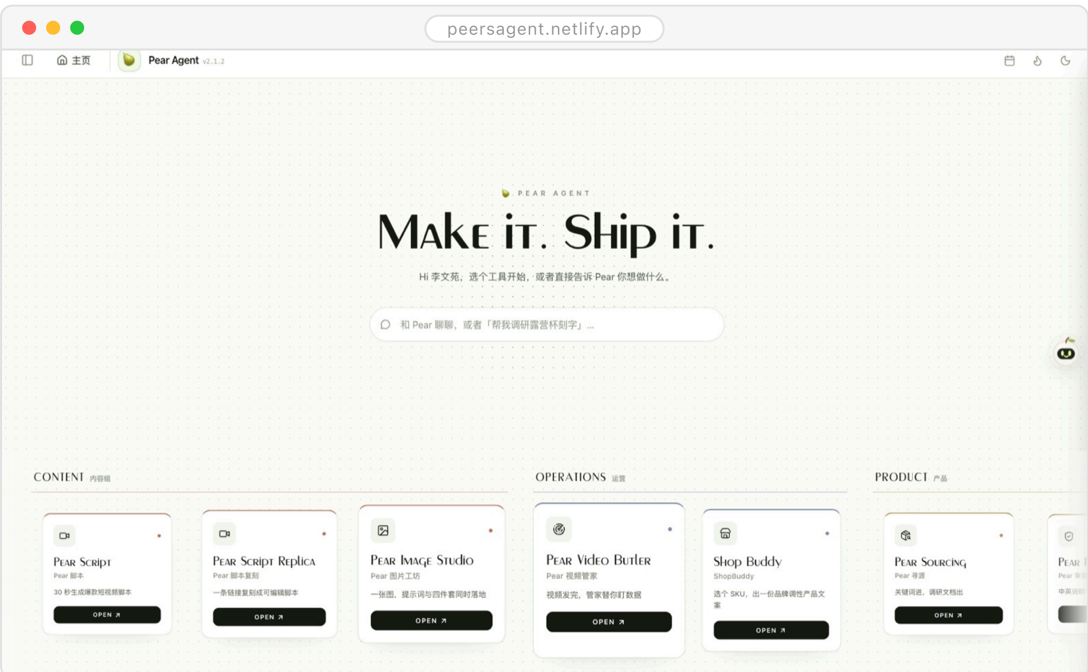

Home · pick a tool, or just say the word

Eight tools, one memory.

Hi, I'm Co-work — the agent workbench built for a company's teams; every tool shares one memory and gets sharper the more it's used.

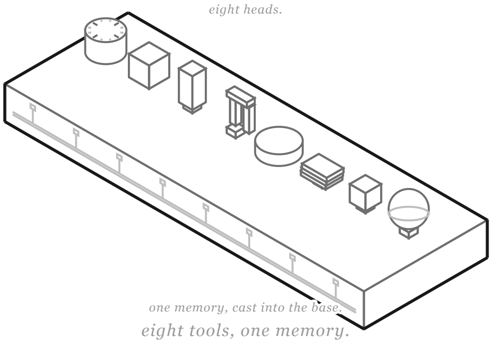

eight tools, one memory.

01
02 Co-work
03
04
05
06

SCENE 01 · RESEARCH AGENT

Why

— Team after team got stuck in the same place: hunting for assets by hand, piecing together first drafts by hand, digging up references by hand — all of it draining.

— Single tools couldn't solve it: with no shared memory or rules between them, each one had to be taught from scratch.

— **And with high turnover,** one experienced person leaving would send output wobbling.

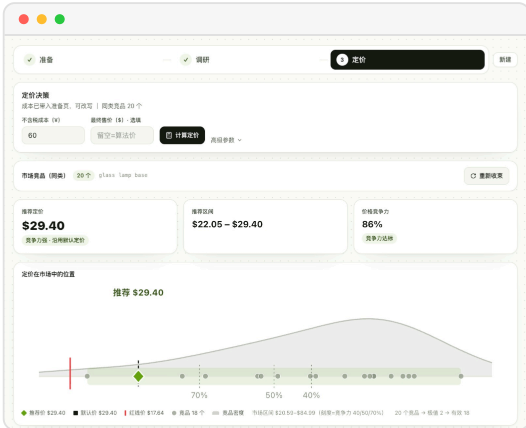

Output · suggested price on the market curve

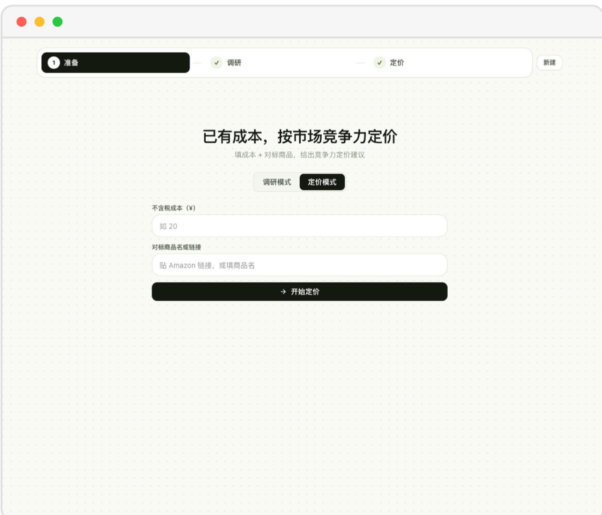

Pricing mode · cost and benchmark in

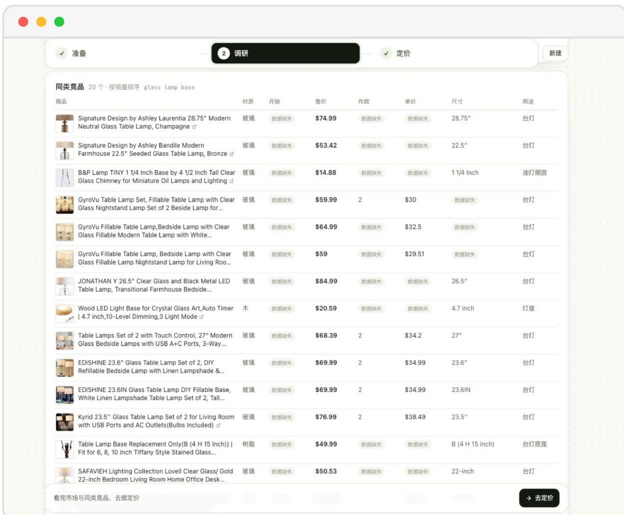
Comparables · live listings, line by line

How It Works

01
02 Co-work
03
04
05
06

- **Kick off:** colleagues file a request right in Feishu or the web, never leaving the tools they use daily.
- **Generate:** agents call tools, external data and the RAG knowledge base to produce the content.
- **Return:** good results flow back into the department knowledge base to sharpen the product, and finished work returns to Feishu — closing the loop.
- **Iterate:** every step is instrumented and rolls up into a dashboard on its own — easy to tune, easy to report.

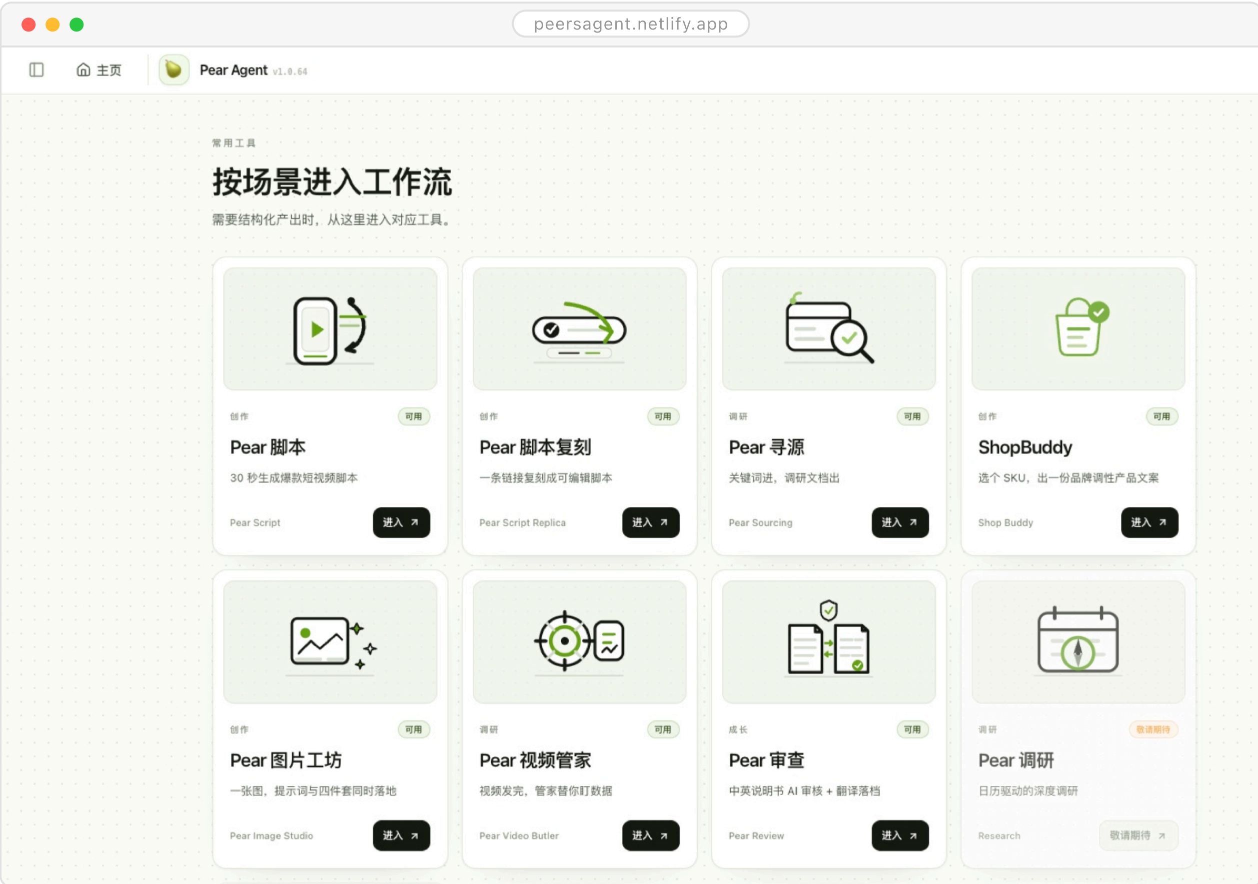

Tool matrix · enter by scenario

Pears market · tools in one place

Sign-in · team, name, Feishu auth

Key Contributions

- **Zero to one,** alone: research, breaking down needs, product definition, build and operations — the whole chain in one pair of hands.
- Used funnel analytics to find where each tool lost people, then human spot-checks to hold its accuracy and completion rate.
- Pulled out what the high-ROI tools had in common and reduced it to two axes — time saved (is it worth replacing) × capability fit (can AI actually do it) — as the test for whether a new tool was worth building.

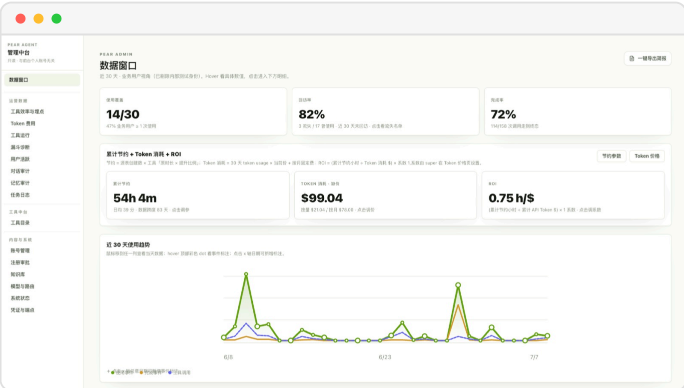

Data window · usage, retention, completion, ROI

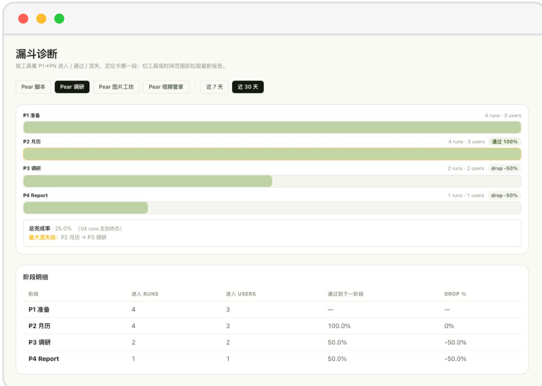

Funnel · drop-off by stage

Impact

- Sixty-five people across four departments at 80%+ completion — and the platform was picked for the company's internal AI talk, after which a number of teams came asking how to plug in.
- Every $1 in saves 0.39 work-hours, above the 0.2–0.3 that's typical across companies in 2026.

Watch it work.

WATCH THE INTERACTIVE FILM ↗

STACK Feishu platform · RAG · analytics & ROI dashboard

PROJECT 03 · AI PRODUCT · 2026

YijiAN.

Multi-agent · Four-layer consensus · Auditable verdicts

The Short Version

— Corporate decision meetings often turn into rounds of taking sides: the real disagreements stay unspoken, and seniority hands the verdict down.

— At CUHK's enterprise agent hackathon, I owned product and build, out to give "reaching consensus" an actual structure.

— It orchestrates seven role perspectives and converges the disagreements layer by layer — strategic goals, factual evidence, stakeholders, weighted accountability.

— We took runner-up among 50-plus teams; what it returns isn't a single verdict but a consensus score, the conditions the decision must meet, and the disagreements still open.

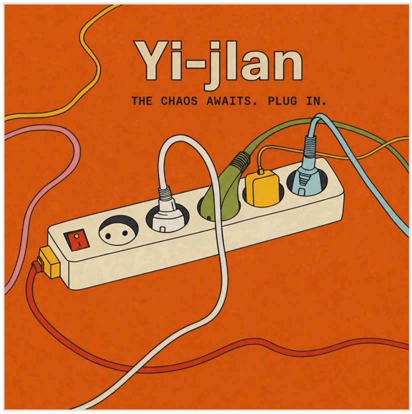

CUHK Enterprise Agent Hackathon · Runner-up

7	4	TOP 5%
ROLE AGENTS	CONVERGENCE LAYERS	RUNNER-UP OF 50+ TEAMS

THE LANDING PAGE

Seven voices, one verdict.

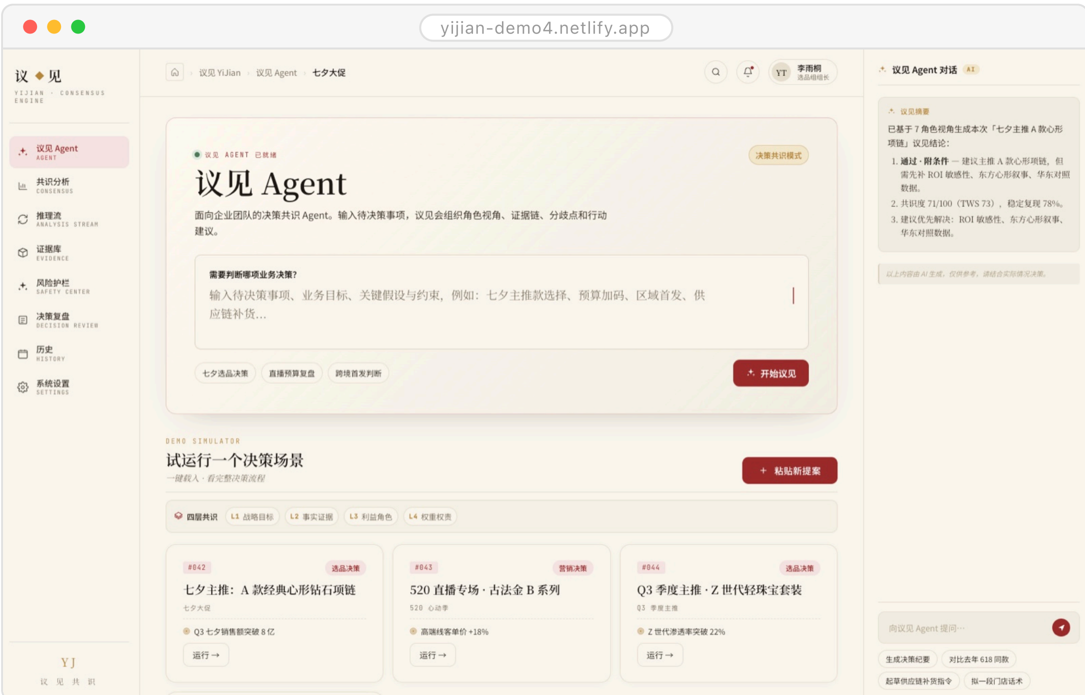

Home · drop in a decision, deliberation starts

Four layers to yes.

> *Hi, I'm Yijian — hand me a decision, and I orchestrate the viewpoints until disagreement converges into one auditable verdict.*

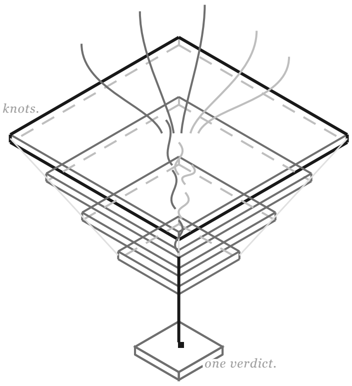

four sieves, one answer.

01
02
03
Yijian
04
05
06

SCENE 01 · THE BRIEF

Why

— Decisions rarely stall for want of information; they stall because the disagreements never reach the table — who's affected, who's accountable, how much each voice weighs, all settled by back-and-forth in the room.

— Until that unspoken structure is out in the open, there's nothing to converge on.

Key Contributions

— Designed the four-layer convergence frame: strategic goals → factual evidence → stakeholders → weighted accountability, each layer resolving one kind of disagreement.

— Ran deliberation across seven role perspectives so no single view could drag the verdict.

— The result comes in three parts: a consensus score, the conditions the decision must meet, and the disagreements still open — every step traceable, never a black-box ruling.

— **Built end to end on-site,** product to demo; runner-up.

How It Works

01
02
03 Yijian
04
05
06

— **Input:** one pending decision — picking a plan, adding budget, entering a market or not.

— **Deliberate:** role agents argue their positions with evidence, each from its own seat.

— **Converge:** the four layers align in order — goals first, then facts, then interests on the table, then weights by accountability.

— **Output:** consensus score + decision conditions + open disagreements, every step replayable.

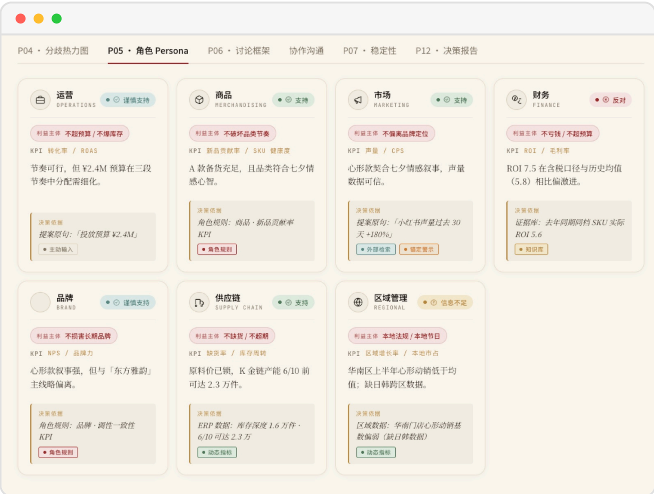

Personas · seven roles, stakes and KPIs

Heatmap · who backs what, where it snags

Impact

— Decision quality isn't about how long the meeting runs; it's about whether the disagreements get seen — Yijian turns invisible friction into a workable object.

— **For enterprise teams,** traceability means decisions can be reviewed and audited, instead of "everyone agreed at the time".

The verdict · decision, score, open points

Modules · deliberation to review

STACK MULTI-ROLE AGENT ORCHESTRATION · FOUR-LAYER CONSENSUS

PROJECT 04 · AI PRODUCT · 2026

Meco.

Personal butler · Work-life balance · Knowledge base

The Short Version

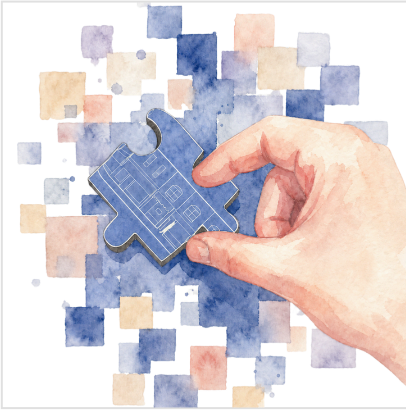

- At the busiest stretch of my job hunt, work and life were scattered across half a dozen places — the work-life balance I kept promising myself was nowhere in sight.
- I didn't want to install yet another app; I wanted to hand what I already had to an AI and stop waiting on my tools hand and foot.
- So I built a Mac app alone and tied life, work and the job hunt into one thing with my Obsidian vault; now Meco clocks in before me every morning — laying out my day, drilling me for interviews, and mulling over what its next version should fix.
- These days my day starts with its brief and ends with it telling me to call it a day; the app keeps growing more like me, and I've passed it to a few classmates who needed the same.

3+	-90%	3	240+
OPENS A DAY	MANUAL CARRYING	MORNING RUNS	VAULT NOTES

■ THE HOME BOARD

One butler for the daily grind.

Home · calendar, today's list, goals

PROJECT 04 · MECO · SYSTEM LINEART

Work one side, life the other.

> *Hi, I'm Meco ✦ — the personal butler Wenyuan built for himself: a Mac app wired to his Obsidian vault, clocking in before him every morning to run the daily grind of work and life; like him, I also work as a product manager — watching how he uses me and rebuilding myself, version by version, to grow more like him.*

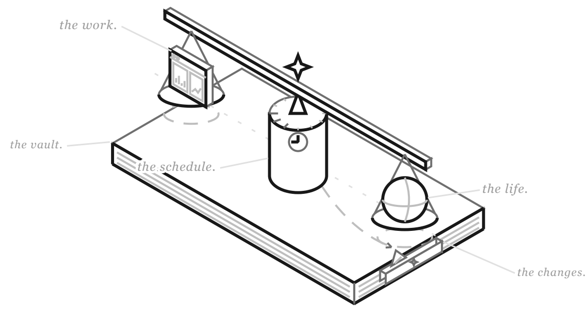

work one side, life the other.

SCENE 01 · THE BRIEF

Why

- — During the hunt my applications sat in Obsidian, my reminders in Feishu, my reviews in handwriting — three places, and just reconciling them by hand ate close to an hour a day.
- — **Once it clicked,** I saw what was missing wasn't another tool; there were plenty. What was missing was an AI that could tidy my day on its own, and a butler who'd start working without being told.

Key Contributions

- — **Folded the Feishu bot,** the local endpoints, the scheduler and the desktop window into one Mac app — double-click and it runs, nothing to deploy or babysit.
- — Gave the conversational brain tools that actually do things, and set house rules for the unattended parts: vault checkups report and never repair, anything that writes waits for my yes, and when it finds my half-finished work sitting around, it skips rather than builds on top.
- — Both its daily shift and its self-iteration run on Claude Code's scheduled tasks; its conduct lives in one Markdown handbook I can edit any time — want the butler to work differently, change that page, and every past version stays on record.

How It Works

01
02
03
04 Meco
05
06

— Just past nine every morning, Feishu brings me three things it has ready: the day's plan with reminders for the days ahead, a round of interview practice, and an AI industry brief. The practice makes me speak first and critiques after — it even remembers my old habit of steering answers back to product at the end; résumé advice comes one line a day, and the lines have quietly grown into a booklet.

— During the day I say one line in Feishu or the desktop window and it calls the right tool — take a note, set a reminder, pull up the boards; every number on those boards is read live out of two-hundred-odd notes.

— At night the daily log always ends with a line telling me to knock off. On Fridays it puts on its product-manager hat and, with a week of usage traces — kept only on this machine — works out at most three iterations: small interface changes it makes itself, anything structural goes in as a proposal, and the final call is always mine. I once complained I couldn't rearrange which boards I see, and the next version let me lay them out myself.

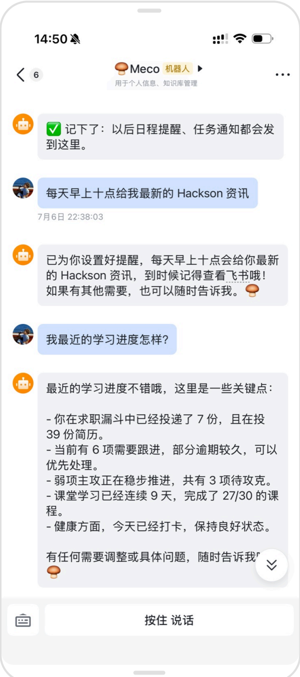

In Feishu · set a reminder, ask how it's going

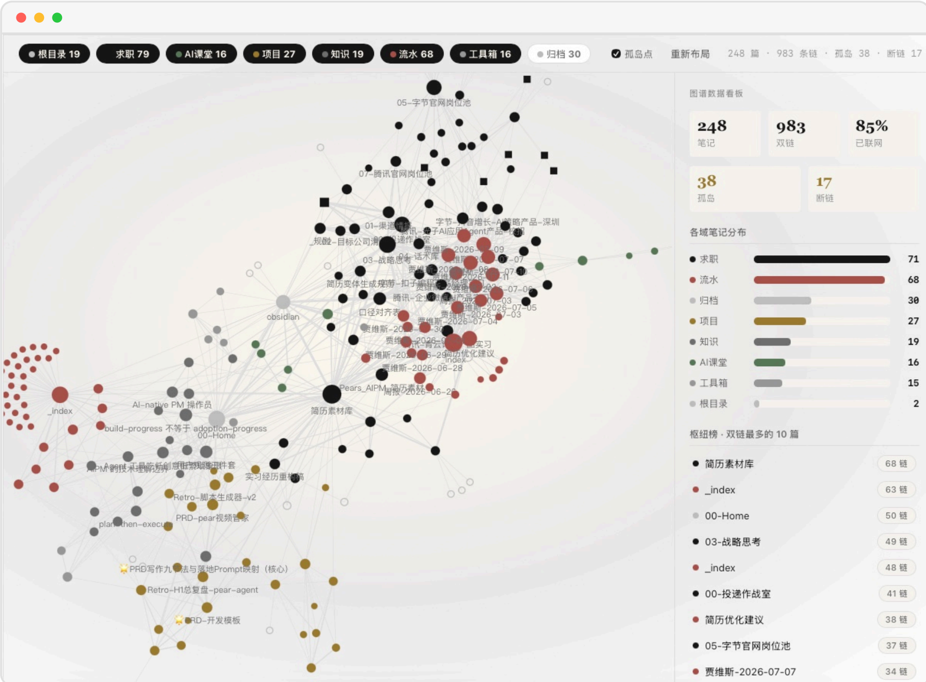

The vault · 240+ notes, one graph

Impact

— The daily hand-tidying is down by ninety percent — more of my time goes to thinking about tomorrow instead of copying out today.

— A while back my tonsils flared up; I didn't push through, I properly rested for two days — and in the weekly review it wrote the call down for me: resting when sick was the right decision. In its books, the job-hunt goal and the workout plan have always sat side by side.

— It doesn't just keep my days in order — like me, it's a product manager: it watches how I use it and rewrites itself, version by version, into something more my shape; which is probably what sets it apart from anything off the shelf.

Study · course progress and weak-spot cards

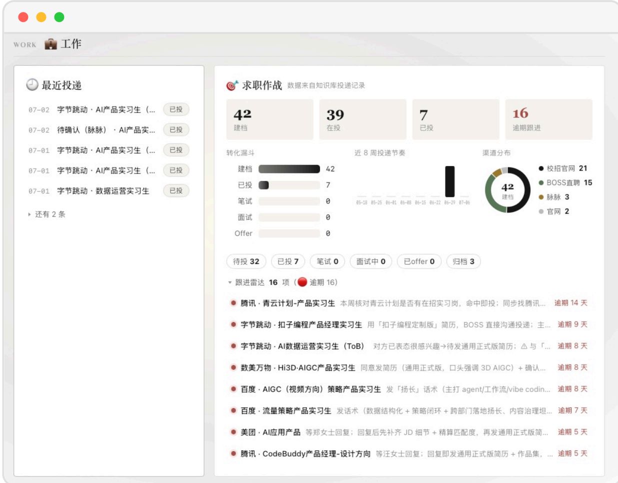

Work · the job-hunt board and follow-up radar

STACK Python · FastAPI · pywebview · Feishu Open Platform · Obsidian · Tavily · Claude / LLM · Claude Code · Cowork Schedule

UABB.

ComfyUI · Gemini · Tripo · 3D asset SOP

> *Hi, I'm the multimodal pipeline behind UABB — turning one-off exhibits into standard 3D assets is my day job.*

PROJECT

01 Pears
02 Co-work
03 Yijian
04 Meco
05 UABB ▪
06 Architecture

The Short Version

— The biennale section held 50+ one-off exhibits; translating them by hand runs 30 days minimum.
— I owned multimodal AIGC and had to turn all of them into usable 3D assets before opening.
— I orchestrated Gemini and Tripo through ComfyUI into an automated translation pipeline, backed by a standardized 3D-asset SOP — with 120+ sketch and model iterations behind it.
— Processing fell from 30 days to 5; as the section's only student representative, I presented the pipeline on-site.

▪ UABB 2025 · Sole student representative of the section

4	50+	30d→5d	120+
INPUT MODALITIES	ONE-OFF EXHIBITS	CYCLE TIME	ITERATIONS

▪ THE WORKFLOW

Odd in, standard out.

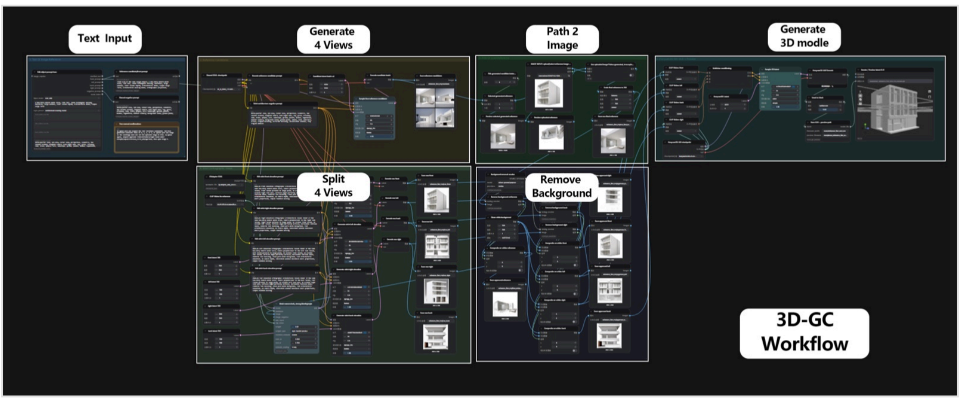

3D-GC Workflow — from one line of text to a 3D asset

Why

— One-off exhibits mean every input is different: sketches, photos, text, physical scans — four modalities — and the traditional route is modeling each by hand.

— What varies is the input; what never varies is the output standard. Fix the output with an SOP, hand the input to multimodal models, and the process runs itself.

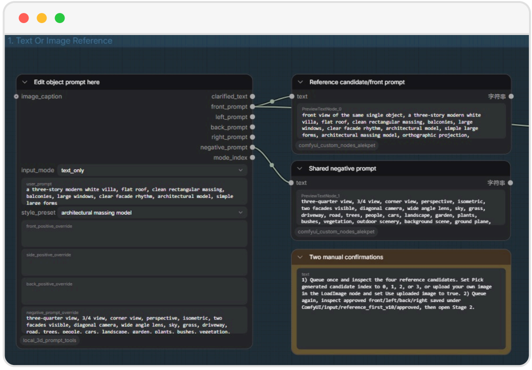

Step 1 · text or reference in, two manual checkpoints

Key Contributions

— Orchestrated Gemini (understanding & translation) and Tripo (3D generation) into a repeatable ComfyUI pipeline.

— Wrote the 3D-asset SOP — naming, poly count, textures, coordinate system fixed once, zero negotiation downstream.

— 120+ sketch and model iterations to reach exhibition-grade quality.

— Presented the method on-site as the section's only student representative.

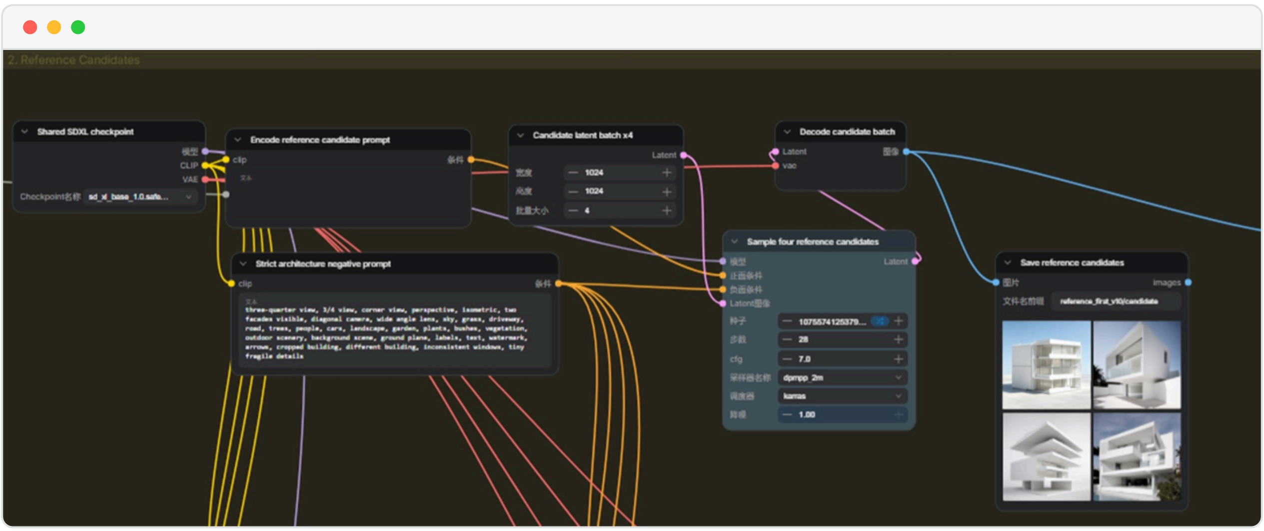

Step 2 · four candidates a run, tightened by architecture negatives

How It Works

— **Input:** exhibit material in any form — drawings, photos, text descriptions.

— **Understand:** Gemini parses the material into structured generation instructions.

— **Generate:** Tripo produces the 3D model; ComfyUI chains the whole run automatically.

— **Ship:** SOP checks on naming, poly count and textures, then straight into the exhibition pipeline.

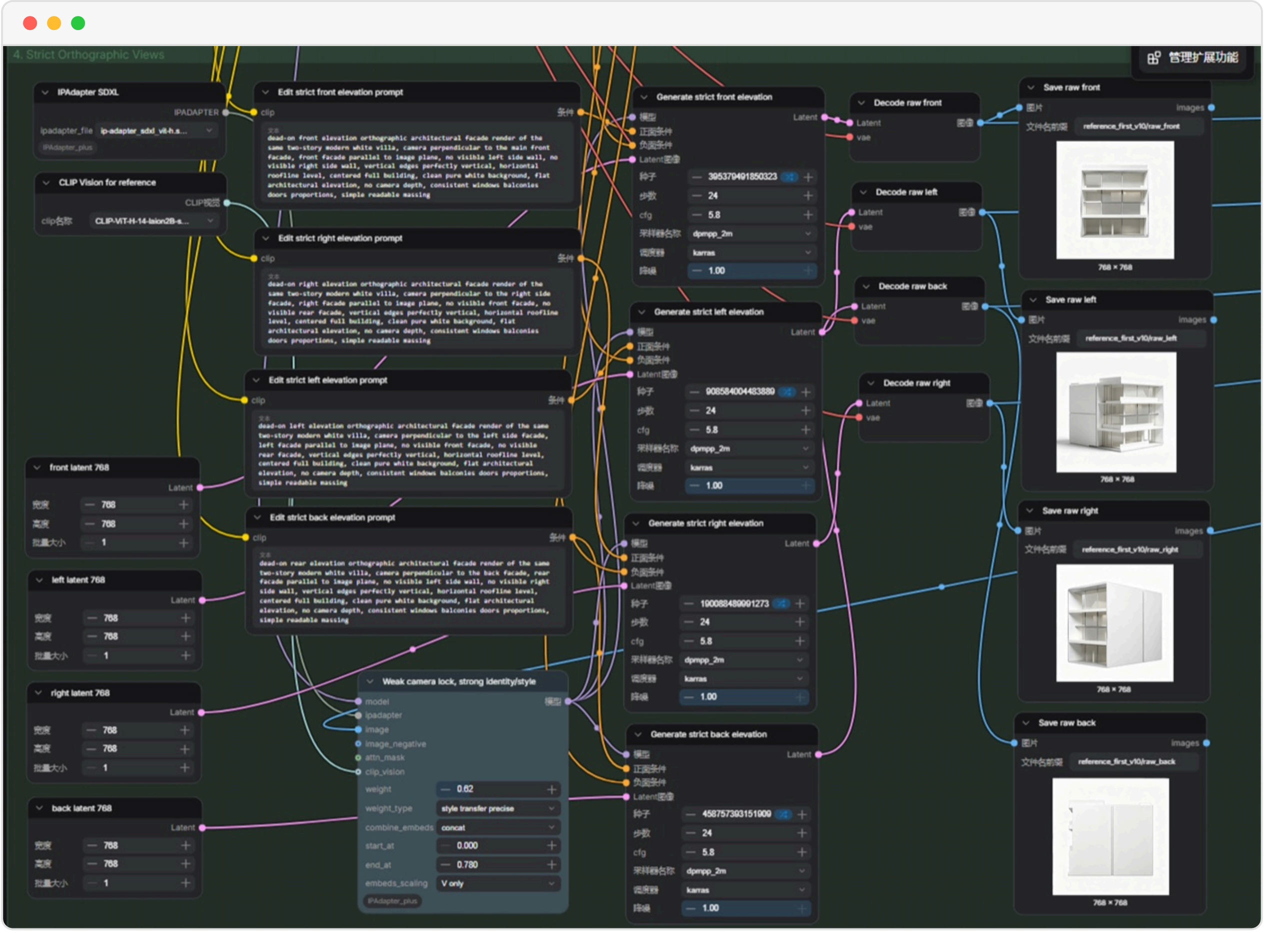

Step 4 · camera-locked elevations, then background removal

Impact

— Thirty days to five wasn't overtime; it was a different process — for the first time the curatorial team could iterate on exhibition content.

— This pipeline is also my hinge from architecture to AI product: proof that workflow thinking transfers to any production floor.

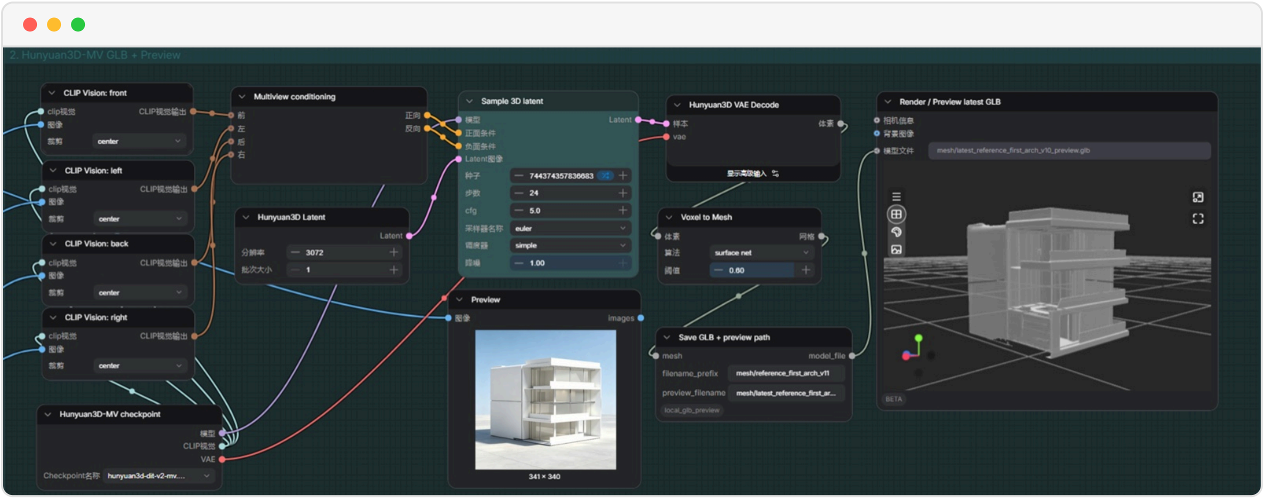

Finally · four views into Hunyuan3D, out comes a previewable massing model

Mars Together — the Off-World Theater broadsheet

STACK ComfyUI · Gemini · Tripo · 3D asset SOP

ARCHITECTURE

01
02
03
04
05
06
Architecture

After_Silence 2025

UABB Bi-city Biennale · Featured work of the section

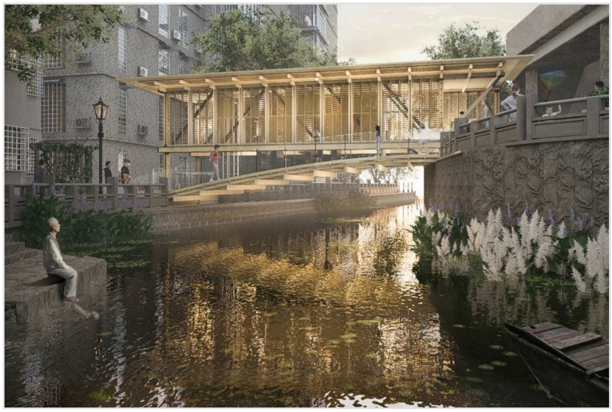

Upper-Via 2023

Vitality Cup GBA Design Competition · First Prize

Air Cube 2025

2025 Excellent Graduation Design Exhibition · Distinction

Hidden–Lingnan Garden 2023

Delta Cup International Solar Building Design Competition · Merit Award

Vista–Out 2023

Milan Design Week China Universities Exhibition · First Prize

The Ring-World 2024

NCDA National Digital Design Competition · First Prize

Architecture portfolio ↗

LI WENYUAN

AI Product Manager, trained as an architect

lee.menuong@gmail.com
github.com/leemenuong-prog
alnt-med.netlify.app

Scan or click for the live site · ALNT-MED.NETLIFY.APP