

AI PRODUCT × ARCHITECTURE

P O R T F O L I O .

AI Product Manager





I'm Wenyan Li (Alint Med). I've shipped AI products before and can take an AI agent from zero to launch on my own — at work I run an agent platform, and I build products of my own. Architecture taught me to turn scattered needs into a stable system; AI product work is where that method meets real scenarios, from idea to launch.

SITE alnt-med.netlify.app




MULTIMODAL Tripo 3D · ComfyUI · Jimeng · Midjourney · UE5

INDEX.

AI PRODUCT × ARCHITECTURE

Pears


AI Product · 2026



01

Co-work


AI Platform · 2026



02

Yi-Jian


AI Product · 2026



03

Meco

AI Product · 2026



04

UABB

AIGC Pipeline · 2025



05

Architecture

ARCHITECTURE · 2023-2025




06

PEARS. ↗

Private agent platform · Work distillation · Workflow agents

The Short Version

- AI has leapt ahead these past two years, but I kept seeing the same gap: piles of work an agent could do were still dragging down people's productivity.
- I wanted to move the bar for building an agent from "explain the requirement to an AI" to "just do the task once and let it watch."
- **Solo**, I built a private-domain agent platform: it watches you work, distills the trace into a PRD you can read and edit, then has AI code it into a workflow agent of your own — zero to launch.
- In the end it took 2nd runner-up at the Tsinghua AI Agent collegiate hackathon, and drew some of the most enterprise interest on site.



■ Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up

TOP 3%

TSINGHUA · 2ND
RUNNER-UP

4+

ENTERPRISE LEADS
ON SITE


36/85

PEER VOTE · NO.1

0→1

SOLO · ZERO TO
LAUNCH

HAND THE BUSYWORK TO A PEAR.



Home · one brief — learning starts

PROJECT

01 Pears ■

02 Co-work

03 Yi-Jian


04 Meco

05 UABB

06 Architecture

WATCH, distill, build.■

Hi, I'm Pears 🍏 — show me
your work once, and I turn it
into a workflow agent of your
own.



SCENE 01 · THE BRIEF

Why


- AI coding brought the bar for writing software way down, yet the front-line colleagues who understand the business best still don't have an "AI employee" of their own.
- **The catch:** the people who know the business aren't necessarily good at describing it — ask them to write prompts or draw flowcharts and you've just raised a new barrier.
- But everyone can do the job they do every day. Treat "doing it once" as the spec, and the barrier is gone.

Key Contributions

- Completed the full run from product definition to launch, solo.
- Designed the observe → distill → generate path: the trace never turns straight into code; it first lands as a PRD you can read and edit, and how far the automation goes is your call.
- Privacy held from day one: it records only when you say so — stop the task, and recording ends on the spot.
- 2nd runner-up at the live roadshow, among the most enterprise-courted projects there (4+ companies), and first in the contestants' own popular vote (36/85 teams).

How It Works

- **Learning:** a short clarifying chat and a goal doc get it clear on what this task should produce.
- **Demo:** you do the work once in the browser as usual, and the extension records the key moves — what you clicked, switched to, typed.
- **Rebuild:** the trace isn't replayed verbatim; RAG reworks it into a workflow that actually runs.
- **Share:** finished capabilities go into an app market — no teaching from scratch, just swap in an app.
- **Boundaries:** it records only after you explicitly start a task, cuts off the moment you stop, and logs only the key events.



① Watch · do it once, key moves recorded



② Assemble · the build direction is your call




③ PRD · readable, editable, handed to AI coding

Impact


- Vibe coding starts from what you say; Pears starts from what you do — and the second never asks you to be good at stating requirements.
- The agent backlog one dev team could never clear becomes something every colleague can clear on their own.
- **A week after launch**, Codex shipped Record & Replay too — outside confirmation that the need here is real.

01 Pears

02
03
04
05
06



Dashboard · runs, hours saved, ROI



Market · pull a working app into your space

WATCH THE PITCH. ■

WATCH THE ROADSHOW FILM ↗




STACK BROWSER EXTENSION · PRD DISTILLATION · AI CODING AGENTS

Co-work. ↗

Agent platform · RAG · Dashboard · Feishu

The Short Version

- During my internship I watched the department's output dragged by two things: piles of work an agent could do were still eating people's hours — and worse, know-how left with whoever quit.
- My job was to upgrade that manual work into agent workflows, while keeping everyone's know-how on the platform.
- So I researched team by team, broke down the needs, and built the department's agent platform zero to one: 8 tools across operations, product and content, iterated on analytics.
- It now serves 65 people across 4 departments at 80%+ completion; by the platform's ROI, every \$1 in saves 0.39 work-hours.



65

PEOPLE · 4 DEPTS

8

PRODUCTION TOOLS


80%+

COMPLETION RATE

0.39h

SAVED PER \$1


Click ONCE — THE AGENTS DO THE REST.



Home

Eight tools, ONE MEMORY.


Hi, I'm Co-work — the agent workbench built for a company's teams; colleagues' know-how keeps settling into skills as they use me, so I get sharper with every run.




SCENE 01 · RESEARCH AGENT

Why


- Team after team got stuck in the same place: hunting for assets by hand, piecing together first drafts by hand, digging up references by hand — all of it draining.
- Single tools couldn't solve it: with no shared memory or rules between them, each one had to be taught from scratch.
- **And with high turnover**, one experienced person leaving would send output wobbling.



① The scenario · a market-competitiveness study




② Comparables · multiple e-commerce platforms plugged in



③ The result · a market-competitiveness analysis

How It Works

- **Kick off:** colleagues file a request right in Feishu or the web, never leaving the tools they use daily.
- **Generate:** agents call tools, external data and the RAG knowledge base to produce the content.
- **Return:** good results flow back into the department knowledge base to sharpen the product, and finished work returns to Feishu — closing the loop.
- **Iterate:** every step is instrumented and rolls up into a dashboard on its own — easy to tune, easy to report.



The tool matrix



The department tool market




The sign-in


01
02
Co-work
03
04
05
06

Key Contributions

- **Zero to one**, alone: research, breaking down needs, product definition, build and operations
 - the whole chain in one pair of hands.
- Used funnel analytics to find where each tool lost people, then human spot-checks to hold its accuracy and completion rate.
- Pulled out what the high-ROI tools had in common and reduced it to two axes — time saved (is it worth replacing) × capability fit (can AI actually do it) — as the test for whether a new tool was worth building.



The data window



The funnel

Impact

- Sixty-five people across four departments at 80%+ completion — and the platform was picked for the company's internal AI talk, after which a number of teams came asking how to plug in.
- Every \$1 in saves 0.39 work-hours, above the 0.2–0.3 that's typical across companies in 2026.

WATCH IT WORK. ■

[WATCH THE INTERACTIVE FILM ↗](#)




STACK FEISHU PLATFORM · RAG · ANALYTICS & ROI DASHBOARD

Yi-Jian. ↗

Multi-agent · Four-layer consensus · Auditable verdicts

The Short Version

- Corporate decision meetings often turn into rounds of taking sides: the real disagreements stay unspoken, and seniority hands the verdict down.
- At CUHK's enterprise agent hackathon, I owned product and build, out to give "reaching consensus" an actual structure.
- It orchestrates seven role perspectives and converges the disagreements layer by layer — strategic goals, factual evidence, stakeholders, weighted accountability.
- We took runner-up among 50-plus teams; what it returns isn't a single verdict but a consensus score, the conditions the decision must meet, and the disagreements still open.



■ CUHK Enterprise Agent Hackathon · Runner-up

TOP 5%

RUNNER-UP OF 50+ TEAMS


7

ROLE AGENTS

4

CONVERGENCE LAYERS


SEVEN VOICES, ONE VERDICT.



Home

FOUR LAYERS TO YES. ■

Hi, I'm Yi-Jian — hand me a decision, and I orchestrate the viewpoints until disagreement converges into one auditable verdict.



01

02

03

Yi-Jian

04

05

06

SCENE 01 · THE BRIEF

Why

- Decisions rarely stall for want of information; they stall because the disagreements never reach the table — who's affected, who's accountable, how much each voice weighs, all settled by back-and-forth in the room.
- Until that unspoken structure is out in the open, there's nothing to converge on.

Key Contributions

- Designed the four-layer convergence frame: strategic goals → factual evidence → stakeholders → weighted accountability, each layer resolving one kind of disagreement.
- Ran deliberation across seven role perspectives so no single view could drag the verdict.
- The result comes in three parts: a consensus score, the conditions the decision must meet, and the disagreements still open — every step traceable, never a black-box ruling.
- **Built end to end on-site**, product to demo; runner-up.

How It Works

- **Input:** one pending decision — picking a plan, adding budget, entering a market or not.
- **Deliberate:** role agents argue their positions with evidence, each from its own seat.
- **Converge:** the four layers align in order — goals first, then facts, then interests on the table, then weights by accountability.
- **Output:** consensus score + decision conditions + open disagreements, every step replayable.

01

02

03


Yi-Jian

■

04

05

06



The modules

Agent stances



The disagreement heatmap

The decision report

Impact

- Decision quality isn't about how long the meeting runs; it's about whether the disagreements get seen — Yi-Jian turns invisible friction into a workable object.
- **For enterprise teams,** traceability means decisions can be reviewed and audited, instead of "everyone agreed at the time".


STACK MULTI-ROLE AGENT ORCHESTRATION · FOUR-LAYER CONSENSUS

Meco. ↗

Personal butler · Work-life balance · Knowledge base

The Short Version

- **At my busiest stretch**, work and life were scattered across half a dozen places — the work-life balance I kept promising myself was nowhere in sight.
- I didn't want to install yet another app; I wanted to hand what I already had to an AI and stop waiting on my tools hand and foot.
- So I built a Mac app alone and tied life, work and the job hunt into Meco with my Obsidian vault; now it clocks in before me every morning — laying out my day, walking me through interviews, and mulling over what its own next version should fix.
- These days my day starts with its report and ends with it telling me to knock off; the app keeps growing more like me, and I've shared it with a few classmates who needed the same.



5+

OPENS A DAY

-90%

MANUAL TIDYING

4

BOARDS

240+

VAULT NOTES

ONE BUTLER FOR THE daily GRIND.




The home board

Work one side, life the other. ■

01
02
03
04
05
06

Meco ■

Hi, I'm Meco ♦ — the personal butler Wenyan built for himself: an app wired to his Obsidian vault, running the daily grind of his work and life; like him, I moonlight as a product manager — watching how he uses me and revising myself, version by version, to fit him better.



SCENE 01 · THE BRIEF

Why


- **At the peak**, my applications sat in Obsidian, my reminders in Feishu, my reviews in handwriting — no attention left for life itself, and just reconciling them by hand ate close to an hour a day.
- **Once it clicked**, I saw what was missing wasn't another tool; there were plenty. What was missing was an AI that could tidy my day on its own, and a butler who'd start working without being told.

Key Contributions


- **Folded Feishu**, the vault, the heartbeat engine (Claude) and the local endpoints into one app — double-click and it runs, no fleet of software to babysit.
- Gave the conversational brain tools that actually do things: every day it tidies my life and work information on its own, and proposes where the next version should improve.
- Both its daily shift and its self-iteration run on Claude Code's scheduled tasks; HITL keeps my hand on the wheel of where the app goes.

How It Works

- **At nine every morning**, Feishu brings what it has ready: the day's plan, a round of coursework, an AI brief — all of it accumulating into a knowledge base of my own.
- **Quick call**: I say one line in Feishu and it reads the vault to call the right tool — take a note, set a reminder, pull up the boards.
- **Iterate**: I gave it an AI-product identity of its own; from my usage data it plans the next version's updates, the same way I would.




In Feishu




The vault

Impact

- The daily hand-tidying is down by ninety percent — more of my time goes to thinking about tomorrow instead of tracing over today.
- It doesn't just keep my days in order — like me, it's a product manager: it watches how I use it and rewrites itself, version by version, into something more my shape; which is probably what sets it apart from anything off the shelf.



The study board



The work board

STACK PYTHON · FASTAPI · FEISHU OPEN PLATFORM · CLAUDE CODE

UABB.

ComfyUI · 3D asset workflow

The Short Version

PROJECT

01 Pears

02 Co-work

03 Yi-Jian

04 Meco

05 UABB ■

06 Architecture

- The biennale section held 50+ one-off exhibits; traditional hand translation runs 30 days minimum.
- I owned the multimodal AIGC workflow, and had to turn all of them into usable glasses-free 3D assets before opening.
- I built an automated translation workflow — ComfyUI orchestrating external AI models — and set up a standardized 3D-asset SOP.
- Processing fell from 30 days to 5; as the section's only student representative, I presented the workflow on-site.



■ UABB 2025 · Sole student representative of the section

4

INPUT MODALITIES

50+

ONE-OFF EXHIBITS


30d → 5d

CYCLE TIME

120+

ITERATIONS


Odd in, STANDARD OUT.



3D-GC Workflow — from one line of text to a 3D asset

Why


- One-off exhibits mean every input is different: sketches, photos, text, physical scans — four modalities.
- The traditional route models each piece by hand, and leaves no room for quick revisions.



Step 1 · text in (or an image, or a model)

Key Contributions


- Orchestrated AI models and 3D generation into a repeatable ComfyUI pipeline.
- Wrote the standardized 3D-asset SOP — naming, textures, coordinate system fixed once, zero merge cost downstream.
- 120+ sketch and model iterations, polishing rough massing into exhibition-grade quality.




Step 2 · pick from the candidates

How It Works

- **Input:** exhibit material in any form — images, models, text descriptions.
- **Generate:** Hunyuan produces the 3D model; ComfyUI chains the whole run automatically.
- **Accept:** SOP checks on naming and textures, then straight into the exhibition candidates.



Step 4 · camera-locked elevations, then background removal



The final result



Mars Together — the Off-World Theater broadsheet

ARCHITECTURE


01
02
03
04
05
06
Architecture



AFTER_SILENCE

2025

UABB Bi-city Biennale · Featured work of the section



UPPER-VIA

2023

Vitality Cup GBA Design Competition · First Prize



AIR CUBE

2025

2025 Excellent Graduation Design Exhibition · Distinction



HIDDEN-LINGNAN GARDEN

2023

Delta Cup International Solar Building Design Competition · Merit Award



VISTA-OUT

2023

Milan Design Week China Universities Exhibition · First Prize



THE RING-WORLD

2024

NCDA National Digital Design Competition · First Prize

[Architecture portfolio ↗](#)



LI WENYUAN

AI Product Manager

lee.menuong@gmail.com

github.com/leemenuong-prog

alnt-med.netlify.app



Scan or click for the live site · ALNT-MED.NETLIFY.APP