

AI PRODUCT

P O R T F O L I O .

AI Product Manager, trained as an architect





Hello, I'm WENYUAN.

AI Product Manager, trained as an architect

I'm Wen Yuan Li (Alnt Med). Trained as an architect, now building AI products — I run a department agent platform at XTOOL and build my own. Architecture taught me to turn scattered needs into a stable system; AI product work is where that method meets real scenarios.

EMAIL lee.menuong@gmail.com

GITHUB github.com/leemenuong-prog

SITE alnt-med.netlify.app



alnt-med.netlify.app

Scan or click for the live

site ↗

EDUCATION

2025 – Shenzhen University · M.Arch, architecture × AI — the pivot to AI product

2020–2025 Guangzhou University · B.Arch

EXPERIENCE

2026 XTOOL · AIPM intern · Co-work agent platform, zero to one (65 people across 4 departments, 80%+ completion)

2025 Puwei Tech · AI product intern (vertical RAG agents)

2025 UABB Bi-city Biennale · multimodal AIGC lead · sole student representative

2025 Architectural Journal (T1) · second author, AIGC spatial optimisation

AWARDS

Challenge Cup · Grand Prize TOP 1%

Lanqiao Cup, national software competition · National First Prize TOP 1%

Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up TOP 3%

CUHK Enterprise Agent Hackathon · Runner-up TOP 5%

NCDA National Digital Design Award · First Prize

SKILLS

PRODUCT AI coding (Claude Code / Cursor) · prototyping (Figma / Axure / ClaudeDesign) · PRD

DATA SQL · Python · A/B testing · AI eval sets · Apify · Excel pivots

ORCHESTRATION Coze · n8n


MULTIMODAL Tripo 3D · ComfyUI · Jimeng · Midjourney · UE5

INDEX.

SIX WORKS · AI PRODUCT × ARCHITECTURE

Pears


AI Product · 2026



01

Co-work


AI Platform · 2026



02

Yijian


AI Product · 2026



03

Meco

AI Product · 2026



04

UABB


AIGC Pipeline · 2025



05

Architecture

SIX WORKS OF ARCHITECTURE · 2023-2025




06

PEARS.

Private agent platform · Work distillation · Workflow agents

The Short Version

- AI has leapt ahead these past two years, but I kept seeing the same gap: teams still dragging through repetitive work that an agent could take off their hands.
- I wanted to move the bar for building an agent from "explain the requirement to an AI" to "just do the task once and let it watch."
- **Solo**, I built a private-domain agent platform: it watches you work, distills the trace into a PRD you can read and edit, then hands it to an AI coding agent that builds a workflow agent of your own — zero to launch.
- It took 2nd runner-up at the Tsinghua AI Agent collegiate hackathon, and drew some of the most enterprise interest on site.



■ Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up

4+

ENTERPRISE LEADS
ON SITE

36/85

PEER VOTE · NO.1

0→1

SOLO · ZERO TO
LAUNCH


TOP 3%

TSINGHUA · 2ND
RUNNER-UP

■ THE LANDING PAGE

HAND THE BUSYWORK TO A PEAR.

把重复工作交给一只梨。



Home · one brief, one sample — learning starts

PROJECT

01 Pears ■

02 Co-work

03 Yijian


04 Meco

05 UABB

06 Architecture

WATCH, distill, build.■

Hi, I'm Pears 🍏 — show me
your work once, and I turn it
into a workflow agent of your
own.



SCENE 01 · THE BRIEF

Why

- AI coding brought the bar for writing software way down, yet the front-line colleagues who understand the business best still don't have an "AI employee" of their own.
- **The catch:** the people who know the work are usually the worst at describing it — ask them to write prompts or draw flowcharts and you've just raised a new barrier.
- But everyone can do the job they do every day. Treat "doing it once" as the spec, and the barrier is gone.

Key Contributions

- **Ran the whole thing solo**, product definition to launch — no team to divide the work or fall back on.
- Designed the observe → distill → generate path: the trace never turns straight into code; it first lands as a PRD you can read and edit, and how far the automation goes is your call.
- Privacy was fixed on day one: it records only when you say so, and stops the instant you stop the task.
- 2nd runner-up at the live roadshow, among the most enterprise-courted projects there (4+ companies), and first in the contestants' own popular vote (36/85).

How It Works

01
Pears
02
03
04
05
06


- **Align:** a short clarifying chat and a goal doc get it clear on what this task should produce.
- Demonstrate: you do the work once in the browser as usual, and the extension records the key moves — what you clicked, switched to, typed.
- **Rebuild:** the trace isn't replayed verbatim; RAG reworks it into a workflow that actually runs.
- **Share:** finished capabilities go into an app market — no teaching from scratch, just swap in an app.
- **Boundaries:** it records only after you explicitly start a task, cuts off the moment you stop, and logs only the key events.



① Watch · clicks, switches, keystrokes



② Assemble · each step, your call




③ Build pack · a PRD you can read and edit

Impact


- Vibe coding starts from what you say; Pears starts from what you do — and the second never asks you to be good at stating requirements.
- The agent backlog one dev team could never clear becomes something every colleague can clear on their own.
- **A week after launch**, Codex shipped Record & Replay too — outside confirmation that the need here is real.

01 Pears

02
03
04
05
06



Dashboard · runs, hours saved, ROI



Market · pull a working app into your space

WATCH THE PITCH. ■

WATCH THE ROADSHOW FILM ↗




STACK BROWSER EXTENSION · PRD DISTILLATION · AI CODING AGENTS

Co-work.

Agent platform · RAG · Dashboard · Feishu

The Short Version

- **Interning at XTOOL**, the department's output was dragged by two things: manual repetition and human-run processes — and worse, know-how left with whoever quit.
- My job was to turn that manual work into agent workflows, and keep everyone's know-how in the platform.
- **Team by team**, I built the department's agent platform zero to one: 8 tools across operations, product and content, iterated on analytics.
- It now serves 65 people across 4 departments at 80%+ completion; by the platform's ROI, every \$1 in saves 0.39 work-hours.



65

PEOPLE · 4 DEPTS

8

PRODUCTION TOOLS

80%+

COMPLETION RATE


0.39h

SAVED PER \$1

■ THE LANDING PAGE

Click ONCE — THE AGENTS do THE REST.

点击，然后让 Agent 替你干活。



Home · pick a tool, or just say the word

PROJECT

01 Pears

02 Co-work ■

03 Yijian


04 Meco

05 UABB

06 Architecture

Eight tools, ONE MEMORY.


Hi, I'm Co-work — the agent workbench built for a company's teams; every tool shares one memory and gets sharper the more it's used.




SCENE 01 · RESEARCH AGENT

Why


- Team after team got stuck in the same place: hunting for assets by hand, piecing together first drafts by hand, digging up references by hand — all of it draining.
- Single tools couldn't solve it: with no shared memory or rules between them, each one had to be taught from scratch.
- **And with high turnover**, one experienced person leaving would send output wobbling.



Output · suggested price on the market curve




Pricing mode · cost and benchmark in



Comparables · live listings, line by line

How It Works

- **Kick off:** colleagues file a request right in Feishu or the web, never leaving the tools they use daily.
- **Generate:** agents call tools, external data and the RAG knowledge base to produce the content.
- **Return:** good results flow back into the department knowledge base to sharpen the product, and finished work returns to Feishu — closing the loop.
- **Iterate:** every step is instrumented and rolls up into a dashboard on its own — easy to tune, easy to report.



Tool matrix · enter by scenario



Pears market · tools in one place




Sign-in · team, name, Feishu auth


01
02
Co-work
03
04
05
06

Key Contributions

- **Zero to one**, alone: research, breaking down needs, product definition, build and operations
 - the whole chain in one pair of hands.
- Used funnel analytics to find where each tool lost people, then human spot-checks to hold its accuracy and completion rate.
- Pulled out what the high-ROI tools had in common and reduced it to two axes — time saved (is it worth replacing) × capability fit (can AI actually do it) — as the test for whether a new tool was worth building.



Data window · usage, retention, completion, ROI



Funnel · drop-off by stage

Impact

- Sixty-five people across four departments at 80%+ completion — and the platform was picked for the company's internal AI talk, after which a number of teams came asking how to plug in.
- Every \$1 in saves 0.39 work-hours, above the 0.2–0.3 that's typical across companies in 2026.

WATCH it work. ■

[WATCH THE INTERACTIVE FILM ↗](#)




STACK FEISHU PLATFORM · RAG · ANALYTICS & ROI DASHBOARD

Yijian.

Multi-agent · Four-layer consensus · Auditable verdicts

The Short Version

- Corporate decision meetings often turn into rounds of taking sides: the real disagreements stay unspoken, and seniority hands the verdict down.
- At CUHK's enterprise agent hackathon, I owned product and build, out to give "reaching consensus" an actual structure.
- It orchestrates seven role perspectives and converges the disagreements layer by layer — strategic goals, factual evidence, stakeholders, weighted accountability.
- We took runner-up among 50-plus teams; what it returns isn't a single verdict but a consensus score, the conditions the decision must meet, and the disagreements still open.



■ CUHK Enterprise Agent Hackathon · Runner-up

7

ROLE AGENTS

4

CONVERGENCE LAYERS


TOP 5%

RUNNER-UP OF 50+ TEAMS

■ THE LANDING PAGE

SEVEN VOICES, ONE VERDICT.


七个视角，一个可追溯的结论。



Home · drop in a decision, deliberation starts

FOUR LAYERS TO YES. ■

Hi, I'm Yijian — hand me a decision, and I orchestrate the viewpoints until disagreement converges into one auditable verdict.



01

02

03

Yijian ■

04

05

06

SCENE 01 · THE BRIEF

Why

- Decisions rarely stall for want of information; they stall because the disagreements never reach the table — who's affected, who's accountable, how much each voice weighs, all settled by back-and-forth in the room.
- Until that unspoken structure is out in the open, there's nothing to converge on.

Key Contributions


- Designed the four-layer convergence frame: strategic goals → factual evidence → stakeholders → weighted accountability, each layer resolving one kind of disagreement.
- Ran deliberation across seven role perspectives so no single view could drag the verdict.
- The result comes in three parts: a consensus score, the conditions the decision must meet, and the disagreements still open — every step traceable, never a black-box ruling.
- **Built end to end on-site**, product to demo; runner-up.

How It Works

- **Input:** one pending decision — picking a plan, adding budget, entering a market or not.
- **Deliberate:** role agents argue their positions with evidence, each from its own seat.
- **Converge:** the four layers align in order — goals first, then facts, then interests on the table, then weights by accountability.
- **Output:** consensus score + decision conditions + open disagreements, every step replayable.

01
02
03
Yifan

04
05
06



Personas · seven roles, stakes and KPIs



Heatmap · who backs what, where it snags

Impact

- Decision quality isn't about how long the meeting runs; it's about whether the disagreements get seen — Yijian turns invisible friction into a workable object.
- **For enterprise teams**, traceability means decisions can be reviewed and audited, instead of "everyone agreed at the time".

01
02
03
04
05
06



The verdict · decision, score, open points



Modules · deliberation to review


STACK MULTI-ROLE AGENT ORCHESTRATION · FOUR-LAYER CONSENSUS

Meco.

Personal butler · Self-growth · Knowledge base

The Short Version

- At the busiest stretch of my job hunt, work and life were scattered across half a dozen places — the work-life balance I kept promising myself was nowhere in sight.
- I didn't want to install yet another app; I wanted to hand what I already had to an AI and stop waiting on my tools hand and foot.
- So I built a Mac app alone, using my Obsidian vault to tie life, work and the job hunt into one thing — Meco — with 8 tools I reach for every day.
- Meco is an app I now open three-plus times a day, and I've since passed it to a few classmates who needed the same.



3+

OPENS A DAY

-90%

MANUAL CARRYING

4+

DAILY ROUTINES


200+

VAULT NOTES

THE HOME BOARD

ONE BUTLER FOR THE daily GRIND.

日常琐事，交给管家。



Home · calendar, today's list, goals

PROJECT

01 Pears

02 Co-work

03 Yijian


04 Meco

05 UABB

06 Architecture

THREE LIVES, ONE THREAD.

Hi, I'm Meco 🍄 — the personal butler Wenyuan built for himself: a Mac app that ties his Feishu bot, his AI tools, his schedule and his Obsidian vault together, and runs the daily grind of work and life for him.



SCENE 01 · THE BRIEF

Why


- During the hunt my applications sat in Obsidian, my reminders in Feishu, my reviews in handwriting — three places, and just reconciling them by hand ate close to an hour a day.
- **Once it clicked**, I saw what was missing wasn't another tool; there were plenty. What was missing was an AI that could tidy my day on its own, and a butler willing to do the work for me.

Key Contributions


- **Folded the Feishu bot**, the local endpoints, the scheduler and the desktop window into one Mac app — double-click and it runs, nothing to deploy or babysit.
- Gave the conversational brain tools that actually do things, and drew a read-only boundary around unattended runs — anything that writes waits for my explicit yes.
- **Turned applications**, health, goals and learning into one screen of charts, and wired Claude Code to Feishu — when a long job finishes, it reports back on its own.

How It Works

- I say a line in Feishu or the desktop window; it catches what I mean and calls the right tool — take a note, set a reminder, check the boards, search the web.
- Every number on the boards is parsed live from an Obsidian vault of 200-plus notes, so one look tells me where I stand.
- **Classes**, reviews, journaling — these are really recurring schedules, and I keep more than four of them running daily: they fire on time, push their results back to Feishu, and feed the fresh material back into the vault along the way.



In Feishu · set a reminder, ask how it's going



The vault · 200+ notes, one graph

Impact

- The daily hand-tidying is down by ninety percent — more of my time goes to thinking about tomorrow instead of copying out today.
- My knowledge base flipped: from me waiting on the tools to the tools coming to me.



Boards · every number parsed live from the vault

STACK PYTHON · FASTAPI · PYWEBVIEW · FEISHU OPEN PLATFORM · OBSIDIAN · TAVILY · CLAUDE / LLM · CLAUDE CODE

UABB.

ComfyUI · Gemini · Tripo · 3D asset SOP

Hi, I'm the multimodal pipeline behind UABB — turning one-off exhibits into standard 3D assets is my day job.

PROJECT

01 Pears

02 Co-work

03 Yijian

04 Meco

05 UABB

06 Architecture

The Short Version

- The biennale section held 50+ one-off exhibits; translating them by hand runs 30 days minimum.
- I owned multimodal AIGC and had to turn all of them into usable 3D assets before opening.
- I orchestrated Gemini and Tripo through ComfyUI into an automated translation pipeline, backed by a standardized 3D-asset SOP — with 120+ sketch and model iterations behind it.
- Processing fell from 30 days to 5; as the section's only student representative, I presented the pipeline on-site.



■ UABB 2025 · Sole student representative of the section

4

INPUT MODALITIES

50+

ONE-OFF EXHIBITS

30d → 5d

CYCLE TIME

120+

ITERATIONS

■ THE POSTER

Odd in, STANDARD OUT.

非标进，标准出。



Mars Together — the Off-World Theater broadsheet

Why

- One-off exhibits mean every input is different: sketches, photos, text, physical scans — four modalities — and the traditional route is modeling each by hand.
- What varies is the input; what never varies is the output standard. Fix the output with an SOP, hand the input to multimodal models, and the process runs itself.

Key Contributions

- Orchestrated Gemini (understanding & translation) and Tripo (3D generation) into a repeatable ComfyUI pipeline.
- Wrote the 3D-asset SOP — naming, poly count, textures, coordinate system fixed once, zero negotiation downstream.
- 120+ sketch and model iterations to reach exhibition-grade quality.
- Presented the method on-site as the section's only student representative.

How It Works

- **Input:** exhibit material in any form — drawings, photos, text descriptions.
- **Understand:** Gemini parses the material into structured generation instructions.
- **Generate:** Tripo produces the 3D model; ComfyUI chains the whole run automatically.
- **Ship:** SOP checks on naming, poly count and textures, then straight into the exhibition pipeline.

Impact

- Thirty days to five wasn't overtime; it was a different process — for the first time the curatorial team could iterate on exhibition content.
- This pipeline is also my hinge from architecture to AI product: proof that workflow thinking transfers to any production floor.

STACK COMFYUI · GEMINI · TRIPO · 3D ASSET SOP


ARCHITECTURE

Six works of architecture — where the systems training comes from. Full volume available separately.

01
02
03
04
05
06
Architecture



AFTER_SILENCE 2025
UABB Bi-city Biennale · Featured work of the section



UPPER-VIA 2023
Vitality Cup GBA Design Competition · First Prize



AIR CUBE 2025
2025 Excellent Graduation Design Exhibition · Distinction



HIDDEN-LINGNAN GARDEN 2023
Delta Cup International Solar Building Design Competition · Merit Award



VISTA-OUT 2023
Milan Design Week China Universities Exhibition · First Prize



THE RING-WORLD 2024
NCDA National Digital Design Competition · First Prize

Architecture portfolio (22 pp.) ↗



LI WENYUAN

AI Product Manager, trained as an architect

lee.menuong@gmail.com

github.com/leemenuong-prog

alnt-med.netlify.app



Scan or click for the live site · ALNT-MED.NETLIFY.APP