AI PRODUCT × ARCHITECTURE

PORTFOLIO.

AI Product Manager

Hello, I'm Wenyuan.

AI Product Manager

I'm Wenyuan Li (Alnt Med). I've shipped AI products before and can take an AI agent from zero to launch on my own — at work I run an agent platform, and I build products of my own. Architecture taught me to turn scattered needs into a stable system; AI product work is where that method meets real scenarios, from idea to launch.

EMAIL	lee.menuong@gmail.com
GITHUB	github.com/leemenuong-prog
SITE	alnt-med.netlify.app

EDUCATION

2025 –	Shenzhen University · M.Arch, architecture × AI
2020–2025	Guangzhou University · B.Arch

EXPERIENCE

2026	XTOOL · AI product intern (AI applications / agents)
2025	Puwei Tech · AI product intern (vertical RAG agents)
2025	UABB Bi-city Biennale · multimodal AIGC lead · sole student representative
2025	Architectural Journal (T1) · second author

AWARDS

Challenge Cup · Grand Prize	TOP 1%
Lanqiao Cup, national software competition · National First Prize	TOP 1%
Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up	TOP 3%
CUHK Enterprise Agent Hackathon · Runner-up	TOP 5%
NCDA National Digital Design Competition · First Prize	TOP 5%

SKILLS

PRODUCT	AI coding (Claude Code / Cursor) · prototyping (Figma / Axure / ClaudeDesign) · PRD
DATA	SQL · Python · A/B testing · AI eval sets · Apify · Excel pivots
ORCHESTRATION	Coze · n8n
MULTIMODAL	Tripo 3D · ComfyUI · Jimeng · Midjourney · UE5

alnt-med.netlify.app

Scan or click for the live site ↗

INDEX.

AI PRODUCT × ARCHITECTURE

PROJECT 01 · AI PRODUCT · 2026

Pears. ↗

Private agent platform · Work distillation · Workflow agents

The Short Version

PROJECT

01 Pears ▪

02 Co-work

03 Yi-Jian

04 Meco

05 UABB

06 Architecture

— AI has leapt ahead these past two years, but I kept seeing the same gap: piles of work an agent could do were still dragging down people's productivity.

— I wanted to move the bar for building an agent from "explain the requirement to an AI" to "just do the task once and let it watch."

— **Solo,** I built a private-domain agent platform: it watches you work, distills the trace into a PRD you can read and edit, then has AI code it into a workflow agent of your own — zero to launch.

— In the end it took 2nd runner-up at the Tsinghua AI Agent collegiate hackathon, and drew some of the most enterprise interest on site.

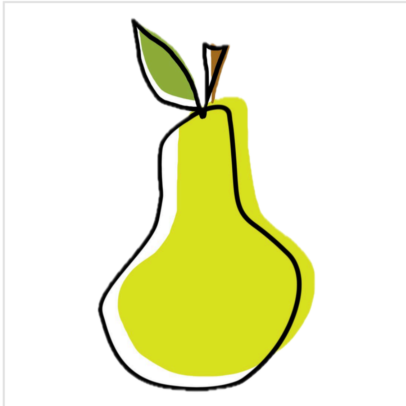

▪ Tsinghua AI Agent Collegiate Hackathon · 2nd runner-up

TOP 3%	4+	36/85	0→1
TSINGHUA · 2ND RUNNER-UP	ENTERPRISE LEADS ON SITE	PEER VOTE · NO.1	SOLO · ZERO TO LAUNCH

Hand the busywork to a pear.

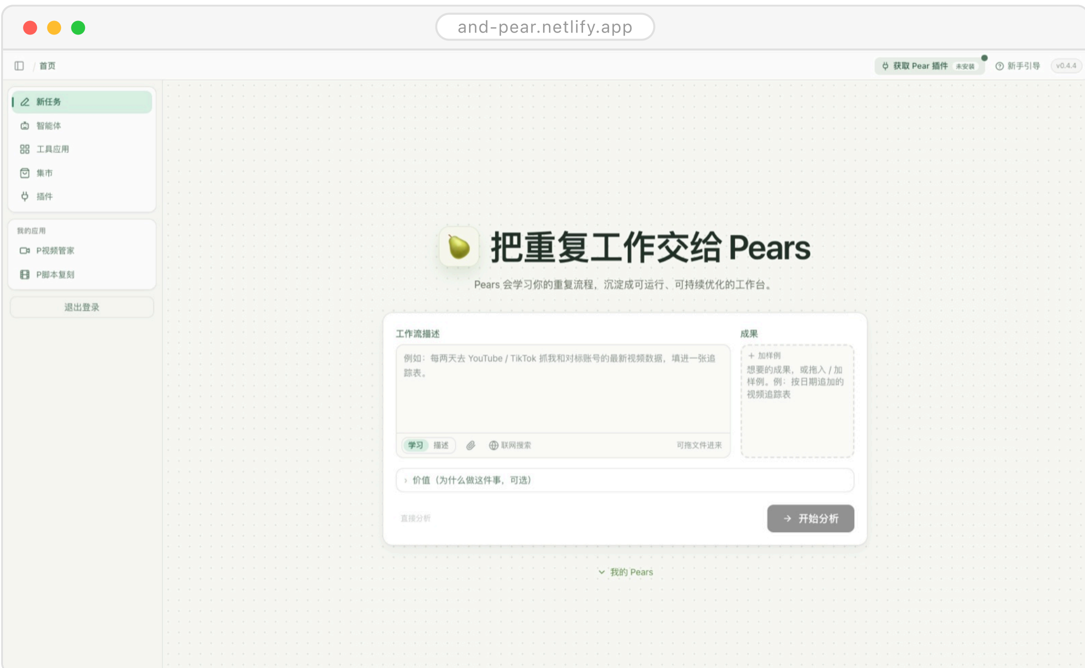

Home · one brief — learning starts

Watch, distill, build.

Hi, I'm Pears 🍐 — show me your work once, and I turn it into a workflow agent of your own.

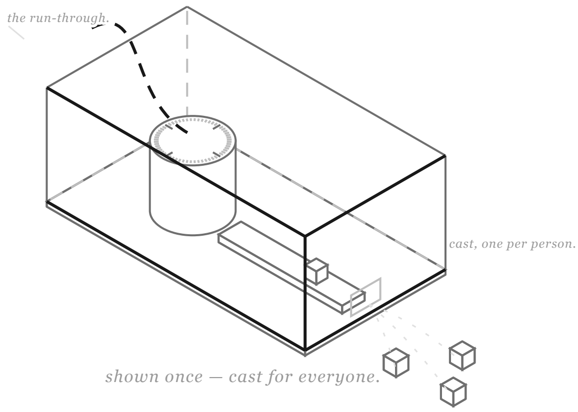

SCENE 01 · THE BRIEF

Why

— AI coding brought the bar for writing software way down, yet the front-line colleagues who understand the business best still don't have an "AI employee" of their own.

— **The catch:** the people who know the business aren't necessarily good at describing it — ask them to write prompts or draw flowcharts and you've just raised a new barrier.

— But everyone can do the job they do every day. Treat "doing it once" as the spec, and the barrier is gone.

Key Contributions

— Completed the full run from product definition to launch, solo.

— Designed the observe → distill → generate path: the trace never turns straight into code; it first lands as a PRD you can read and edit, and how far the automation goes is your call.

— Privacy held from day one: it records only when you say so — stop the task, and recording ends on the spot.

— 2nd runner-up at the live roadshow, among the most enterprise-courted projects there (4+ companies), and first in the contestants' own popular vote (36/85 teams).

How It Works

— **Learning:** a short clarifying chat and a goal doc get it clear on what this task should produce.

— **Demo:** you do the work once in the browser as usual, and the extension records the key moves — what you clicked, switched to, typed.

— **Rebuild:** the trace isn't replayed verbatim; RAG reworks it into a workflow that actually runs.

— **Share:** finished capabilities go into an app market — no teaching from scratch, just swap in an app.

— **Boundaries:** it records only after you explicitly start a task, cuts off the moment you stop, and logs only the key events.

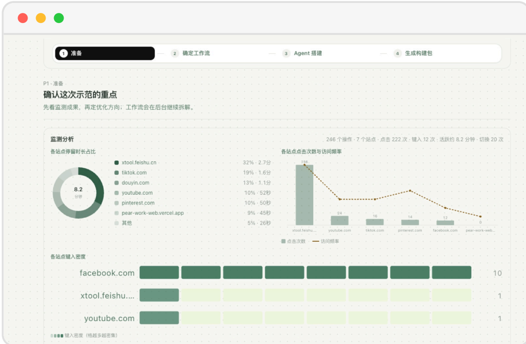

① Watch · do it once, key moves recorded

② Assemble · the build direction is your call

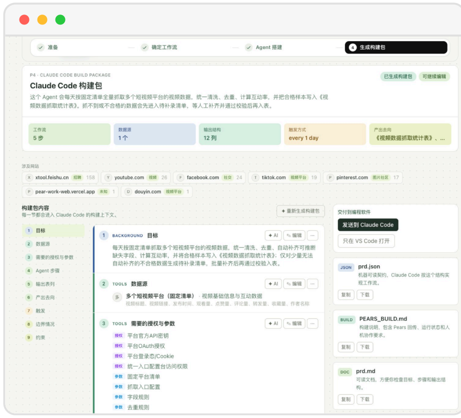

③ PRD · readable, editable, handed to AI coding

Impact

— Vibe coding starts from what you say; Pears starts from what you do — and the second never asks you to be good at stating requirements.

— The agent backlog one dev team could never clear becomes something every colleague can clear on their own.

— **A week after launch,** Codex shipped Record & Replay too — outside confirmation that the need here is real.

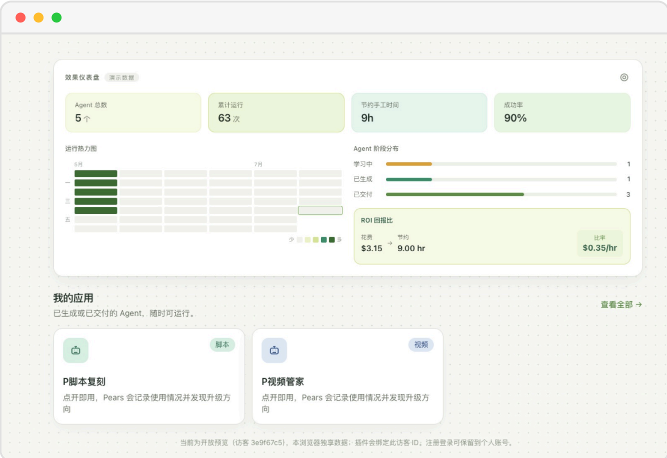

Dashboard · runs, hours saved, ROI

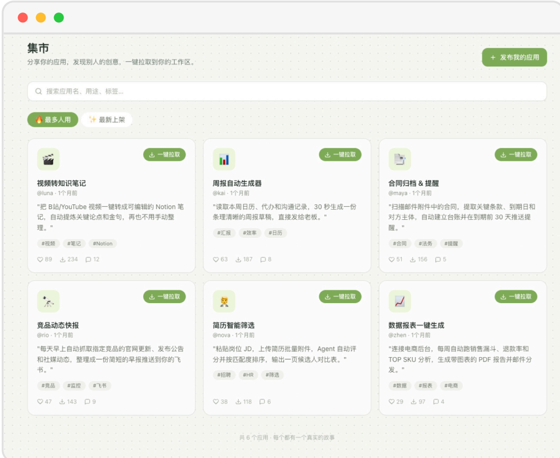

Market · pull a working app into your space

Watch the pitch.

WATCH THE ROADSHOW FILM ↗

STACK Browser extension · PRD distillation · AI coding agents

PROJECT 02 · AI PLATFORM · 2026

Co-work. ↗

Agent platform · RAG · Dashboard · Feishu

The Short Version

- During my internship I watched the department's output dragged by two things: piles of work an agent could do were still eating people's hours — and worse, know-how left with whoever quit.
- My job was to upgrade that manual work into agent workflows, while keeping everyone's know-how on the platform.
- So I researched team by team, broke down the needs, and built the department's agent platform zero to one: 8 tools across operations, product and content, iterated on analytics.
- It now serves 65 people across 4 departments at 80%+ completion; by the platform's ROI, every $1 in saves 0.39 work-hours.

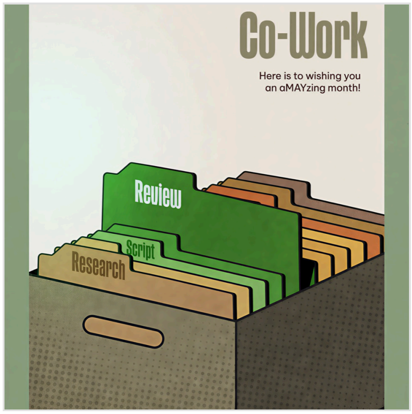

65	8	80%+	0.39h
PEOPLE · 4 DEPTS	PRODUCTION TOOLS	COMPLETION RATE	SAVED PER $1

Click once — the agents do the rest.

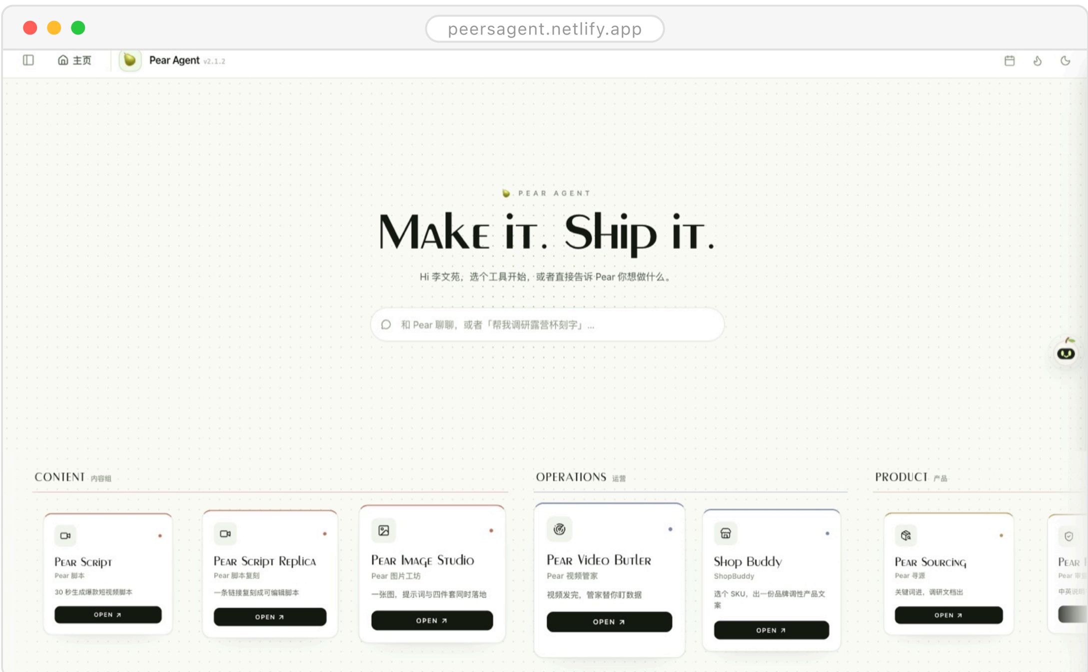

Home

Eight tools, one memory.

01
02
Co-work
03
04
05
06

Hi, I'm Co-work — the agent workbench built for a company's teams; colleagues' know-how keeps settling into skills as they use me, so I get sharper with every run.

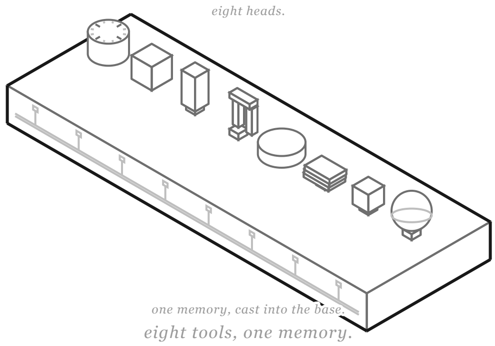

SCENE 01 · RESEARCH AGENT

Why

— Team after team got stuck in the same place: hunting for assets by hand, piecing together first drafts by hand, digging up references by hand — all of it draining.

— Single tools couldn't solve it: with no shared memory or rules between them, each one had to be taught from scratch.

— **And with high turnover,** one experienced person leaving would send output wobbling.

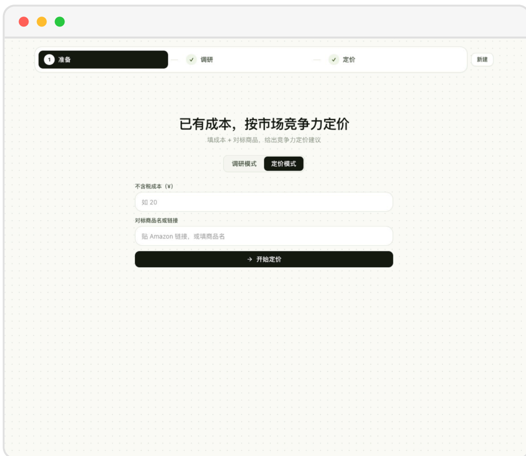

① The scenario · a market-competitiveness study

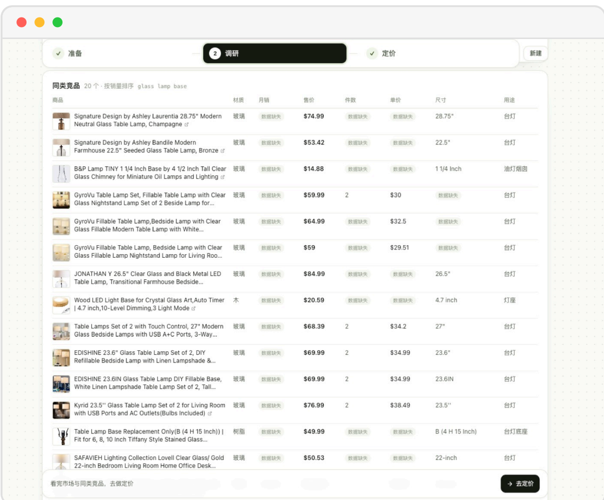

② Comparables · multiple e-commerce platforms plugged in

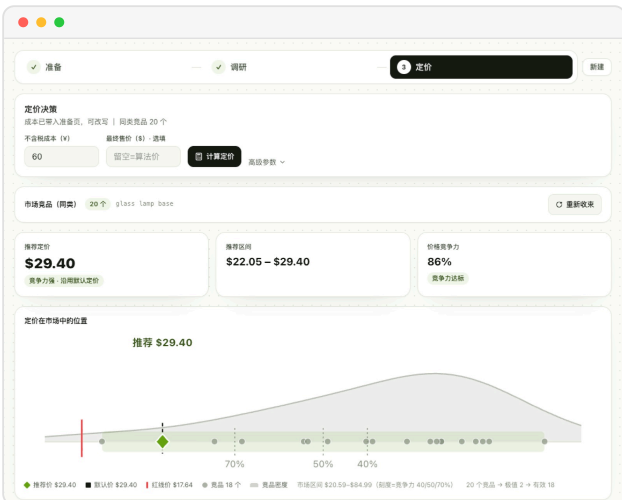

③ The result · a market-competitiveness analysis

How It Works

— **Kick off:** colleagues file a request right in Feishu or the web, never leaving the tools they use daily.

— **Generate:** agents call tools, external data and the RAG knowledge base to produce the content.

— **Return:** good results flow back into the department knowledge base to sharpen the product, and finished work returns to Feishu — closing the loop.

— **Iterate:** every step is instrumented and rolls up into a dashboard on its own — easy to tune, easy to report.

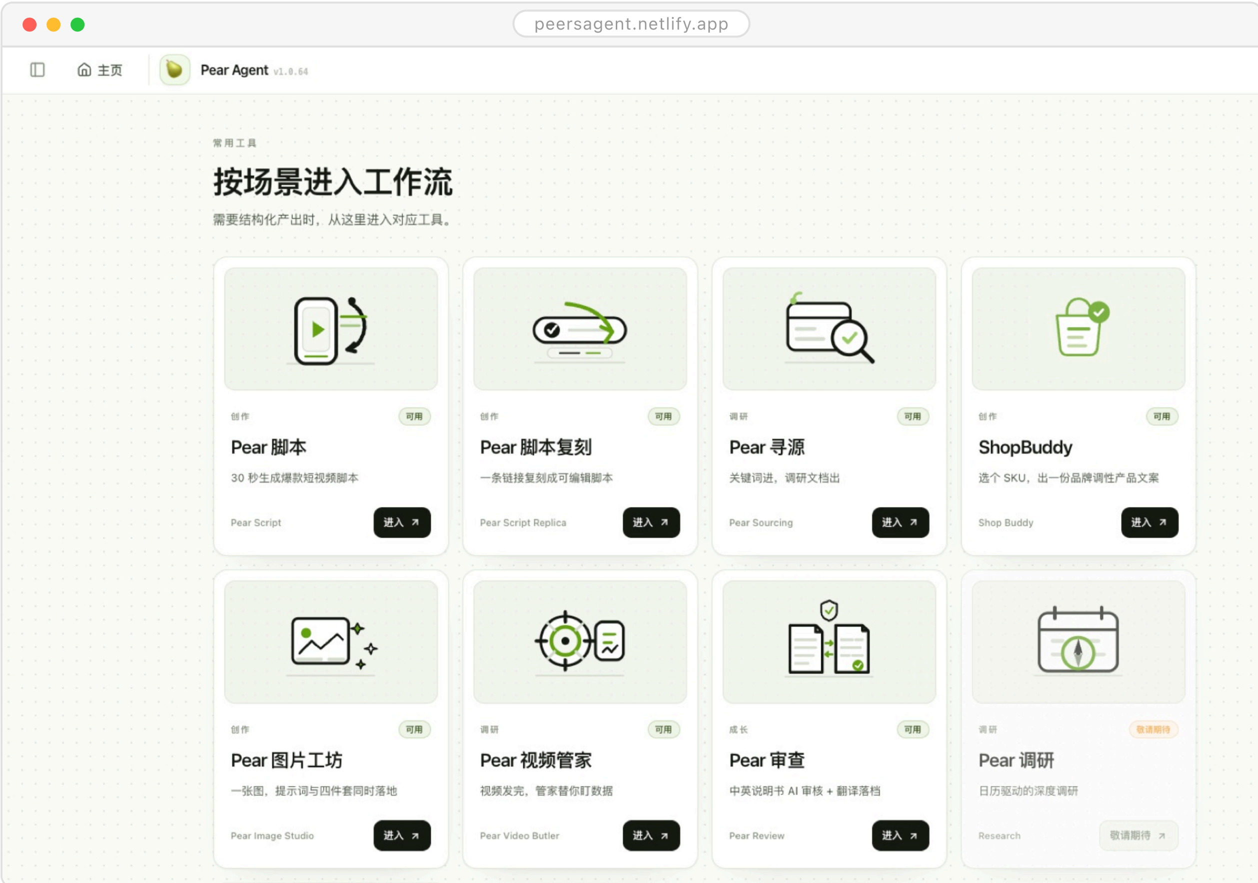

The tool matrix

The department tool market

The sign-in

Key Contributions

01
02 Co-work
03
04
05
06

- **Zero to one,** alone: research, breaking down needs, product definition, build and operations — the whole chain in one pair of hands.
- Used funnel analytics to find where each tool lost people, then human spot-checks to hold its accuracy and completion rate.
- Pulled out what the high-ROI tools had in common and reduced it to two axes — time saved (is it worth replacing) × capability fit (can AI actually do it) — as the test for whether a new tool was worth building.

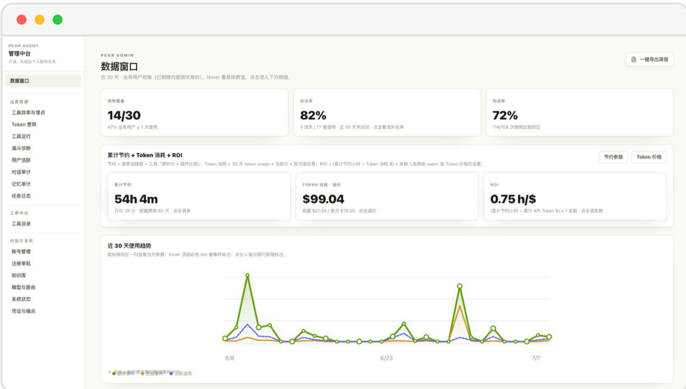

The data window

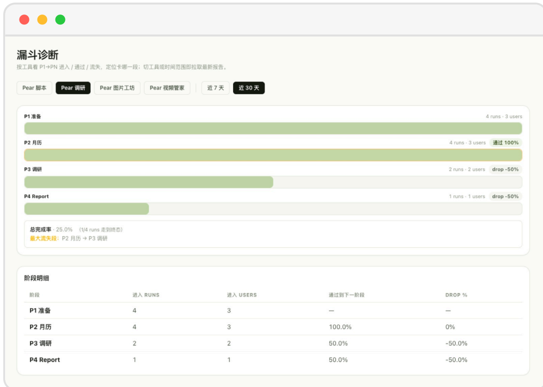

The funnel

Impact

- Sixty-five people across four departments at 80%+ completion — and the platform was picked for the company's internal AI talk, after which a number of teams came asking how to plug in.
- Every $1 in saves 0.39 work-hours, above the 0.2–0.3 that's typical across companies in 2026.

WATCH IT WORK.

WATCH THE INTERACTIVE FILM ↗

STACK FEISHU PLATFORM · RAG · ANALYTICS & ROI DASHBOARD

Yi-Jian. ↗

Multi-agent · Four-layer consensus · Auditable verdicts

The Short Version

— Corporate decision meetings often turn into rounds of taking sides: the real disagreements stay unspoken, and seniority hands the verdict down.

— At CUHK's enterprise agent hackathon, I owned product and build, out to give "reaching consensus" an actual structure.

— It orchestrates seven role perspectives and converges the disagreements layer by layer — strategic goals, factual evidence, stakeholders, weighted accountability.

— We took runner-up among 50-plus teams; what it returns isn't a single verdict but a consensus score, the conditions the decision must meet, and the disagreements still open.

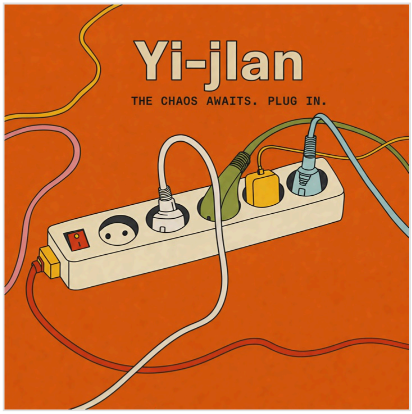

■ CUHK Enterprise Agent Hackathon · Runner-up

TOP 5%	7	4
RUNNER-UP OF 50+ TEAMS	ROLE AGENTS	CONVERGENCE LAYERS

Seven voices, one verdict.

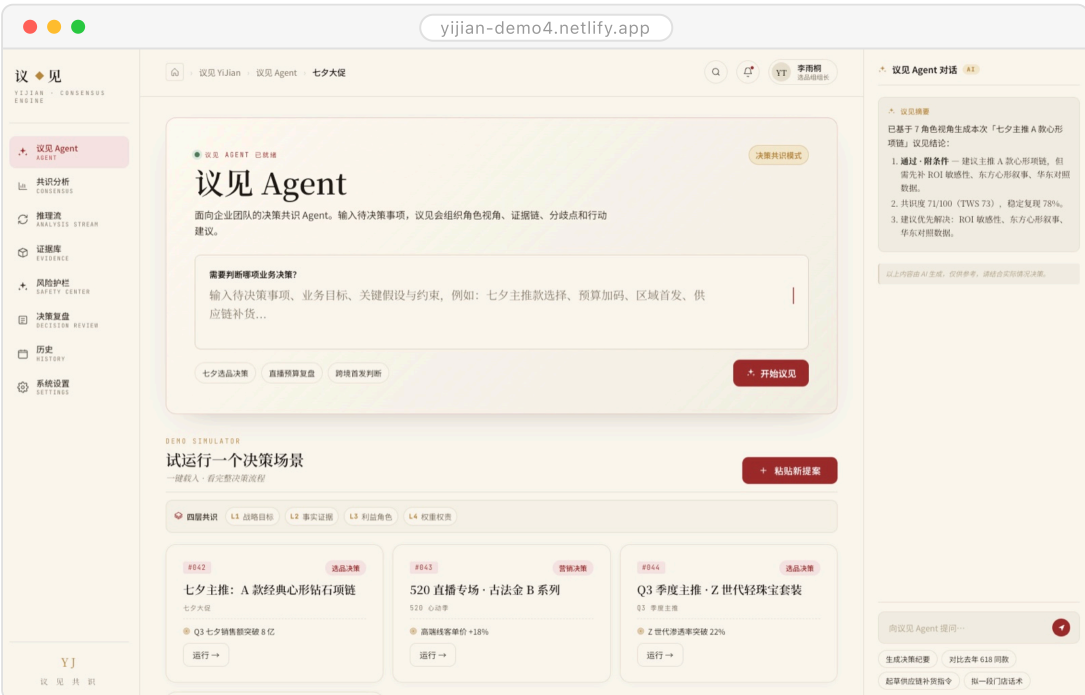

Home

Four layers to yes.

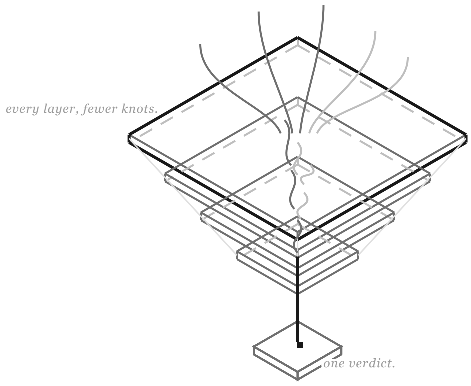

four sieves, one answer.

> *Hi, I'm Yi-Jian — hand me a decision, and I orchestrate the viewpoints until disagreement converges into one auditable verdict.*

01
02
03 Yi-Jian
04
05
06

SCENE 01 · THE BRIEF

Why

— Decisions rarely stall for want of information; they stall because the disagreements never reach the table — who's affected, who's accountable, how much each voice weighs, all settled by back-and-forth in the room.

— Until that unspoken structure is out in the open, there's nothing to converge on.

Key Contributions

— Designed the four-layer convergence frame: strategic goals → factual evidence → stakeholders → weighted accountability, each layer resolving one kind of disagreement.

— Ran deliberation across seven role perspectives so no single view could drag the verdict.

— The result comes in three parts: a consensus score, the conditions the decision must meet, and the disagreements still open — every step traceable, never a black-box ruling.

— **Built end to end on-site,** product to demo; runner-up.

How It Works

— **Input:** one pending decision — picking a plan, adding budget, entering a market or not.

— **Deliberate:** role agents argue their positions with evidence, each from its own seat.

— **Converge:** the four layers align in order — goals first, then facts, then interests on the table, then weights by accountability.

— **Output:** consensus score + decision conditions + open disagreements, every step replayable.

01
02
03 Yi-Jian
04
05
06

The modules

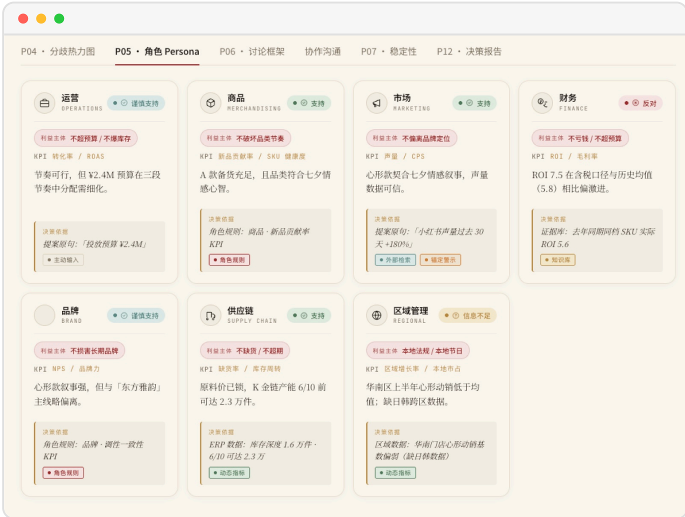

Agent stances

The disagreement heatmap

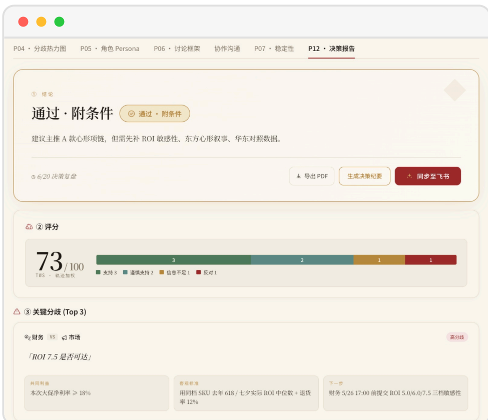

The decision report

Impact

— Decision quality isn't about how long the meeting runs; it's about whether the disagreements get seen — Yi-Jian turns invisible friction into a workable object.

— **For enterprise teams,** traceability means decisions can be reviewed and audited, instead of "everyone agreed at the time".

STACK **Multi-role agent orchestration · four-layer consensus**

PROJECT 04 · AI PRODUCT · 2026

Meco. ↗

Personal butler · Work-life balance · Knowledge base

The Short Version

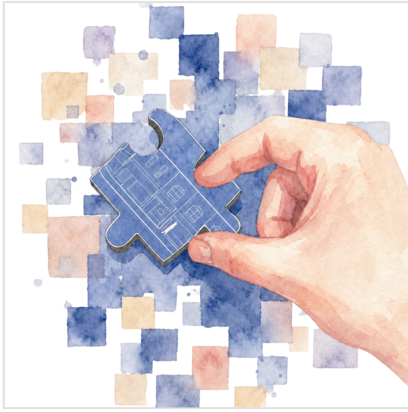

— **At my busiest stretch,** work and life were scattered across half a dozen places — the work-life balance I kept promising myself was nowhere in sight.

— I didn't want to install yet another app; I wanted to hand what I already had to an AI and stop waiting on my tools hand and foot.

— So I built a Mac app alone and tied life, work and the job hunt into Meco with my Obsidian vault; now it clocks in before me every morning — laying out my day, walking me through interviews, and mulling over what its own next version should fix.

— These days my day starts with its report and ends with it telling me to knock off; the app keeps growing more like me, and I've shared it with a few classmates who needed the same.

5+	-90%	4	240+
OPENS A DAY	MANUAL TIDYING	BOARDS	VAULT NOTES

One butler for the daily grind.

The home board

PROJECT 04 · MECO · SYSTEM LINEART

Work one side, life the other.

01
02
03
04 Meco
05
06

> *Hi, I'm Meco ✦ — the personal butler Wenyuan built for himself: an app wired to his Obsidian vault, running the daily grind of his work and life; like him, I moonlight as a product manager — watching how he uses me and revising myself, version by version, to fit him better.*

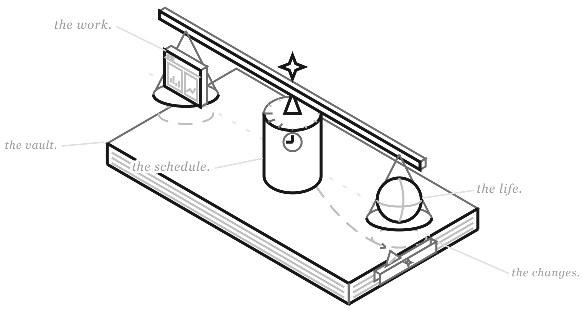

work one side, life the other.

SCENE 01 · THE BRIEF

Why

— **At the peak,** my applications sat in Obsidian, my reminders in Feishu, my reviews in handwriting — no attention left for life itself, and just reconciling them by hand ate close to an hour a day.

— **Once it clicked,** I saw what was missing wasn't another tool; there were plenty. What was missing was an AI that could tidy my day on its own, and a butler who'd start working without being told.

Key Contributions

— **Folded Feishu,** the vault, the heartbeat engine (Claude) and the local endpoints into one app — double-click and it runs, no fleet of software to babysit.

— Gave the conversational brain tools that actually do things: every day it tidies my life and work information on its own, and proposes where the next version should improve.

— Both its daily shift and its self-iteration run on Claude Code's scheduled tasks; HITL keeps my hand on the wheel of where the app goes.

How It Works

01
02
03
04 Meco
05
06

- **At nine every morning,** Feishu brings what it has ready: the day's plan, a round of coursework, an AI brief — all of it accumulating into a knowledge base of my own.
- **Quick call:** I say one line in Feishu and it reads the vault to call the right tool — take a note, set a reminder, pull up the boards.
- **Iterate:** I gave it an AI-product identity of its own; from my usage data it plans the next version's updates, the same way I would.

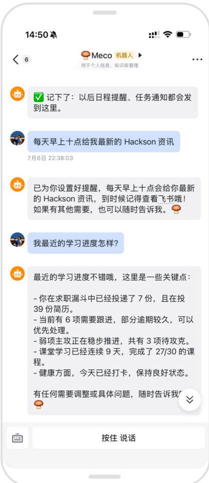

In Feishu

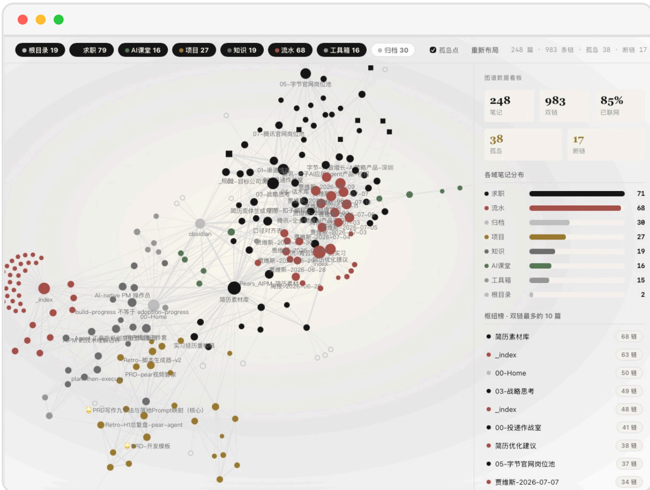

The vault

Impact

- The daily hand-tidying is down by ninety percent — more of my time goes to thinking about tomorrow instead of tracing over today.
- It doesn't just keep my days in order — like me, it's a product manager: it watches how I use it and rewrites itself, version by version, into something more my shape; which is probably what sets it apart from anything off the shelf.

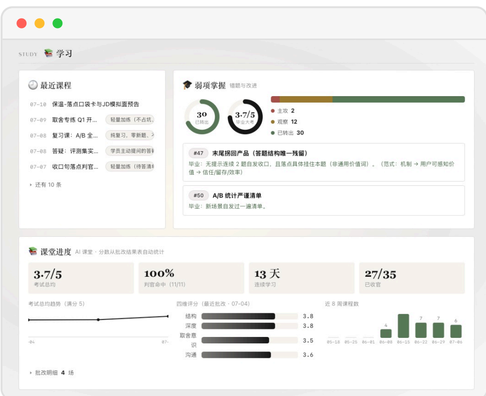

The study board

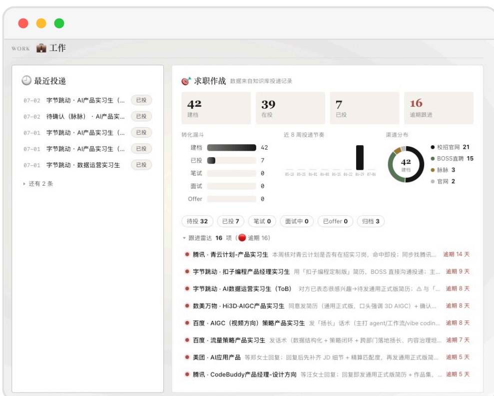

The work board

STACK **Python · FastAPI · Feishu Open Platform · Claude Code**

PROJECT 05 · AIGC PIPELINE · 2025

UABB.

ComfyUI · 3D asset workflow

The Short Version

— The biennale section held 50+ one-off exhibits; traditional hand translation runs 30 days minimum.

— I owned the multimodal AIGC workflow, and had to turn all of them into usable glasses-free 3D assets before opening.

— I built an automated translation workflow — ComfyUI orchestrating external AI models — and set up a standardized 3D-asset SOP.

— Processing fell from 30 days to 5; as the section's only student representative, I presented the workflow on-site.

▪ UABB 2025 · Sole student representative of the section

4	50+	30d→5d	120+
INPUT MODALITIES	ONE-OFF EXHIBITS	CYCLE TIME	ITERATIONS

Odd in, standard out.

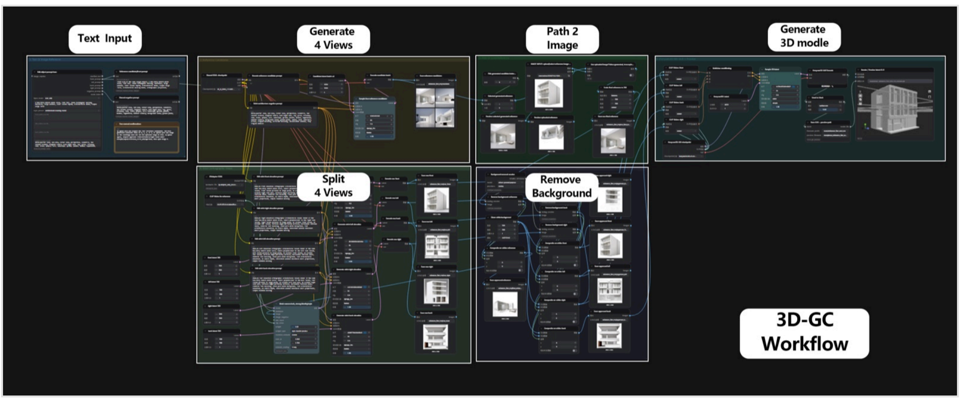

3D-GC Workflow — from one line of text to a 3D asset

SCENE 01 · INPUT & CANDIDATES

Why

01
02
03
04
05
UABB ■
06

— One-off exhibits mean every input is different: sketches, photos, text, physical scans — four modalities.

— The traditional route models each piece by hand, and leaves no room for quick revisions.

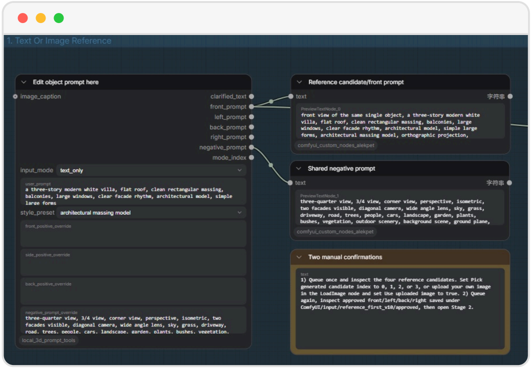

Step 1 · text in (or an image, or a model)

Key Contributions

— Orchestrated AI models and 3D generation into a repeatable ComfyUI pipeline.

— Wrote the standardized 3D-asset SOP — naming, textures, coordinate system fixed once, zero merge cost downstream.

— 120+ sketch and model iterations, polishing rough massing into exhibition-grade quality.

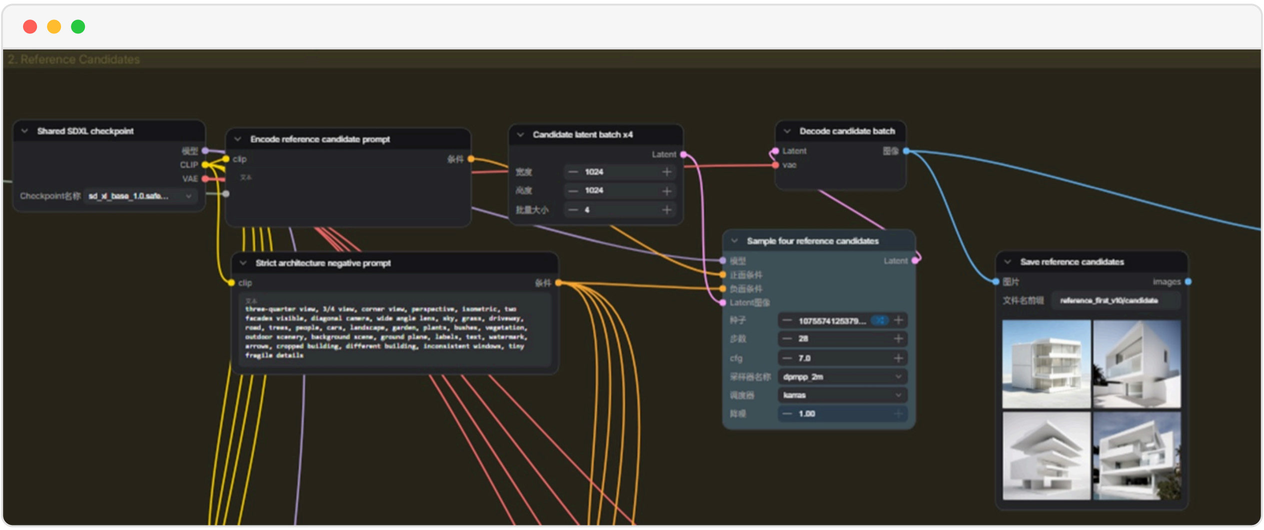

Step 2 · pick from the candidates

How It Works

- **Input:** exhibit material in any form — images, models, text descriptions.
- **Generate:** Hunyuan produces the 3D model; ComfyUI chains the whole run automatically.
- **Accept:** SOP checks on naming and textures, then straight into the exhibition candidates.

01
02
03
04
05 UABB
06

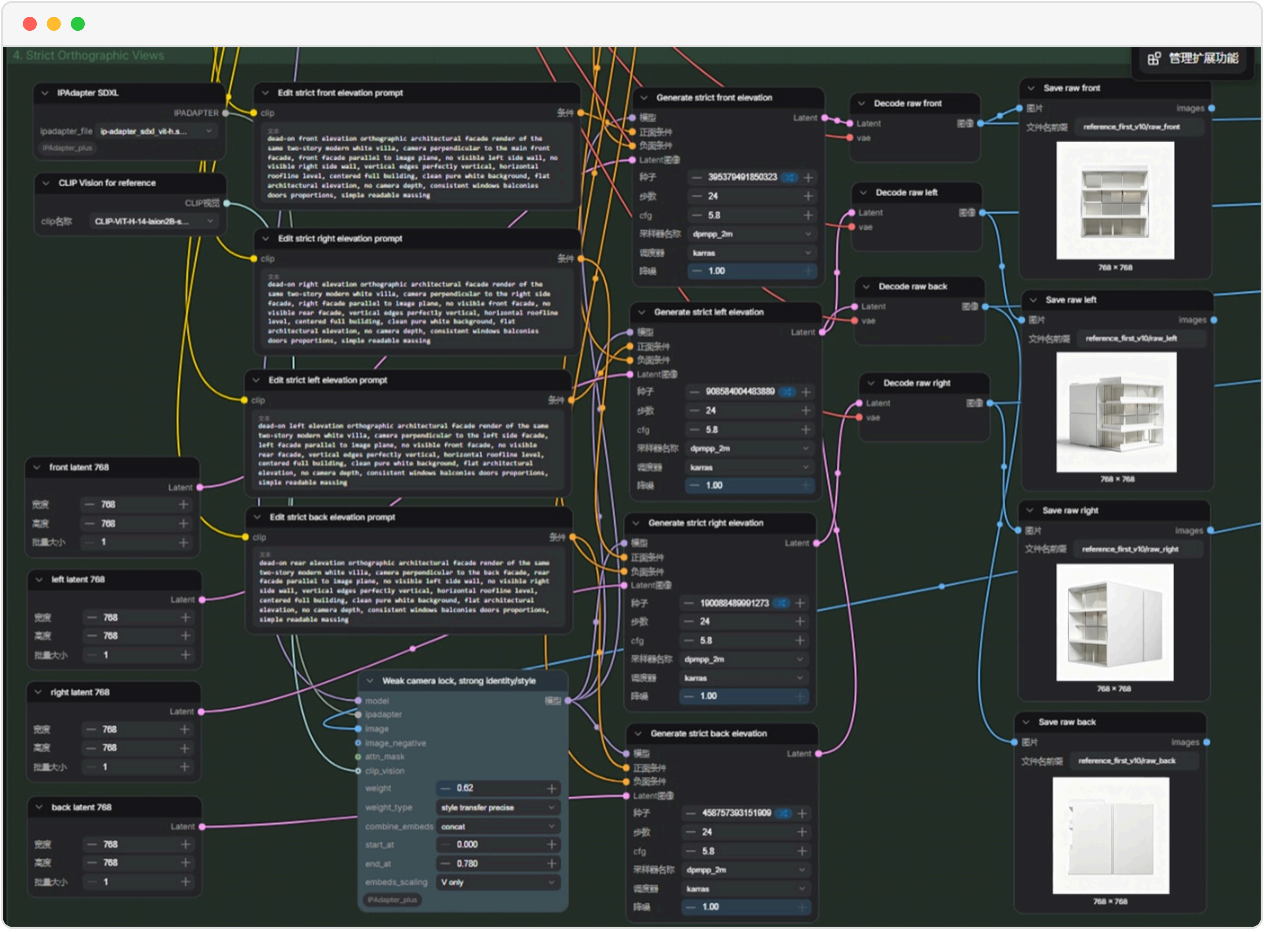

Step 4 · camera-locked elevations, then background removal

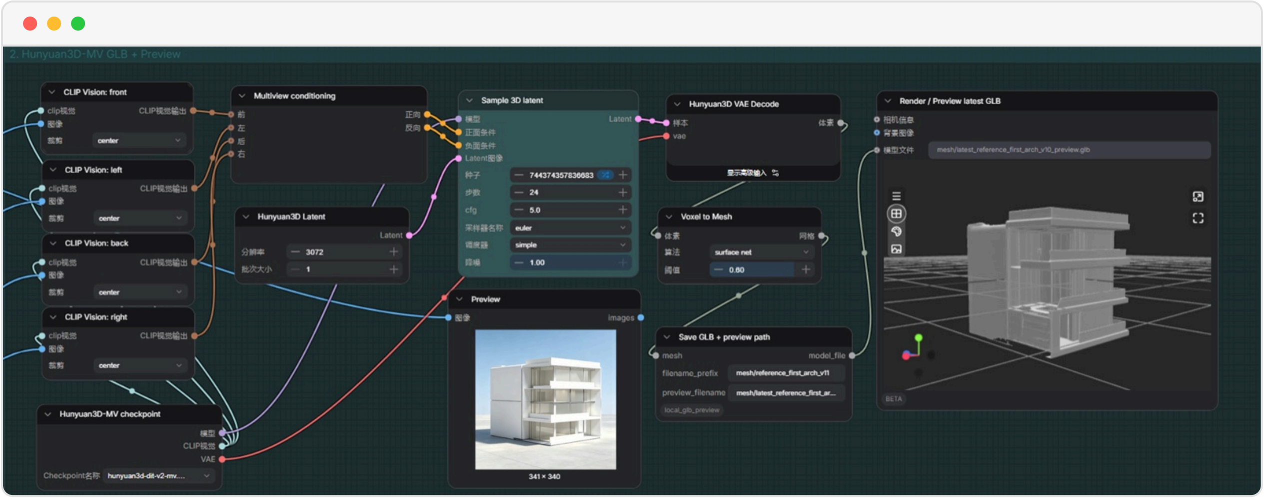

The final result

Mars Together — the Off-World Theater broadsheet

STACK ComfyUI · Hunyuan · 3D asset SOP

ARCHITECTURE

After_Silence 2025

UABB Bi-city Biennale · Featured work of the section

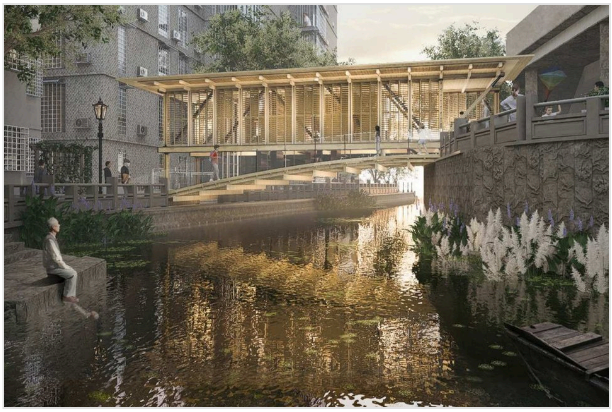

Upper-Via 2023

Vitality Cup GBA Design Competition · First Prize

Air Cube 2025

2025 Excellent Graduation Design Exhibition · Distinction

Hidden–Lingnan Garden 2023

Delta Cup International Solar Building Design Competition · Merit Award

Vista–Out 2023

Milan Design Week China Universities Exhibition · First Prize

The Ring-World 2024

NCDA National Digital Design Competition · First Prize

Architecture portfolio ↗

LI WENYUAN

AI Product Manager

lee.menuong@gmail.com

github.com/leemenuong-prog

alnt-med.netlify.app

Scan or click for the live site · ALNT-MED.NETLIFY.APP